

(Planning with) Dynamic Programming

Introduction

Outline

- ✓ Introduction
- ✓ Dynamic programming
- ✓ Policy Evaluation
- ✓ Policy Iteration
- ✓ Value Iteration
- ✓ Advanced topics
 - ✓ Asynchronous update
 - ✓ Approximated approaches

What is dynamic programming

Dynamic ↪ problem with sequential or temporal component

Programming ↪ optimising a program, i.e. a policy

- ✓ A method for solving complex problems by breaking them down into subproblems
 - ✓ Solve the subproblems
 - ✓ Combine solutions to subproblems
- ✓ It **is not** divide-et-impera
- ✓ Differentiates by **overlapping breakdown**

Requirements for dynamic programming

- ✓ Optimal substructure

- ✓ Principle of optimality applies
 - ✓ Optimal solution can be decomposed into subproblems

- ✓ Overlapping subproblems

- ✓ Subproblems recur many times
 - ✓ Solutions can be cached and reused

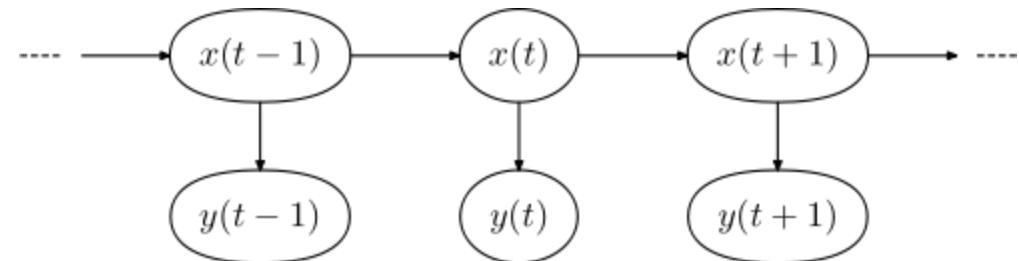
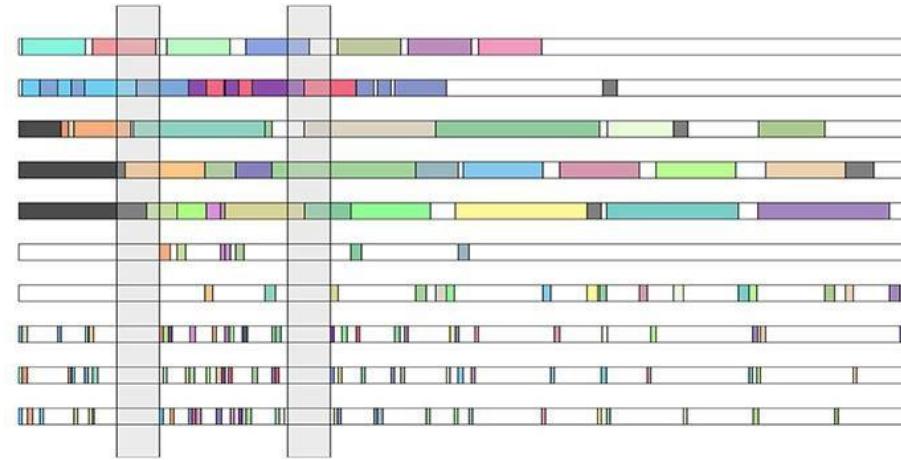
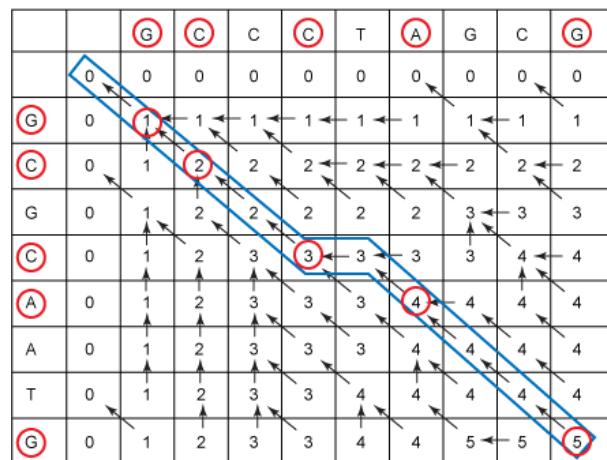
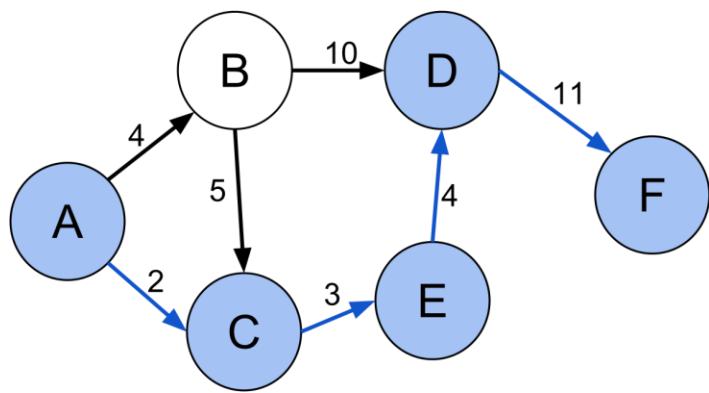
Markov decision processes satisfy both properties

- ✓ Bellman equation gives recursive decomposition
- ✓ Value function stores and reuses solutions

Planning by dynamic programming

- ✓ Dynamic programming assumes full knowledge of the MDP
- ✓ Planning in RL (repetita)
 - ✓ A model of the environment is known
 - ✓ The agent improves its policy
- ✓ Dynamic programming can be used for planning in RL
- ✓ Prediction
 - ✓ **Input:** MDP $\langle \mathcal{S}, \mathcal{A}, \mathbf{P}, \mathcal{R}, \gamma \rangle$ and policy π **or** MRP $\langle \mathcal{S}, \mathbf{P}, \mathcal{R}, \gamma \rangle$
 - ✓ **Output:** value function v_π
- ✓ Control
 - ✓ **Input:** MDP $\langle \mathcal{S}, \mathcal{A}, \mathbf{P}, \mathcal{R}, \gamma \rangle$
 - ✓ **Output:** optimal value function v_{π_*} **and** optimal policy π_*

Applications of Dynamic Programming



Policy Evaluation

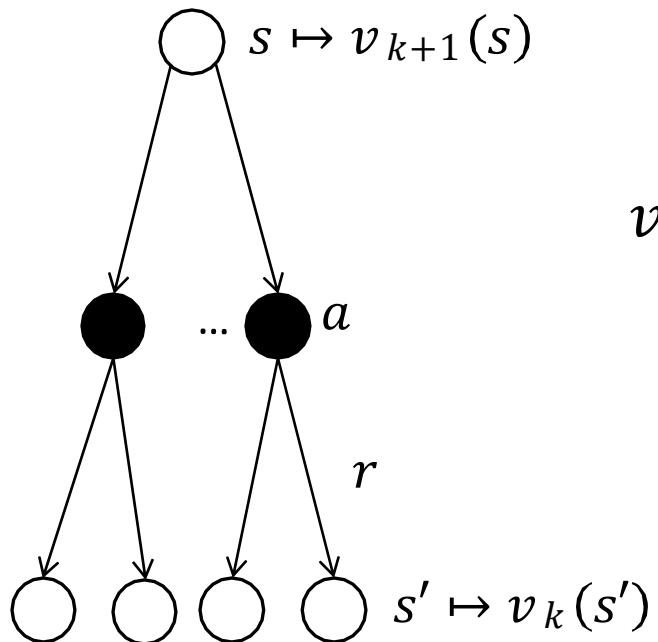
Iterative Policy Evaluation

- ✓ **Problem:** evaluate a given policy π
- ✓ **Solution:** iterative application of Bellman expectation backup

$$v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_\pi$$

- ✓ **Using synchronous backups**
 - i. At each iteration $k + 1$
 - ii. For all states $s \in \mathcal{S}$
 - iii. Update $v_{k+1}(s)$ from $v_k(s')$ where s' is a successor state of s

Iterative Policy Evaluation - Formally



$$v_{k+1}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left(\mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a v_k(s') \right)$$

$$v_{k+1} = \mathcal{R}^\pi + \gamma \mathbf{P}^\pi v_k$$

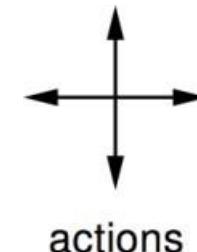
Evaluating a Random Policy in the Small Gridworld

- ✓ Undiscounted episodic MPD ($\gamma = 1$)
- ✓ Nonterminal states $1, \dots, 14$
- ✓ One terminal state (shown twice as shaded squares)
- ✓ Actions leading out of the grid leave state unchanged
- ✓ Reward is -1 until the terminal state is reached
- ✓ Agent follows uniform random policy

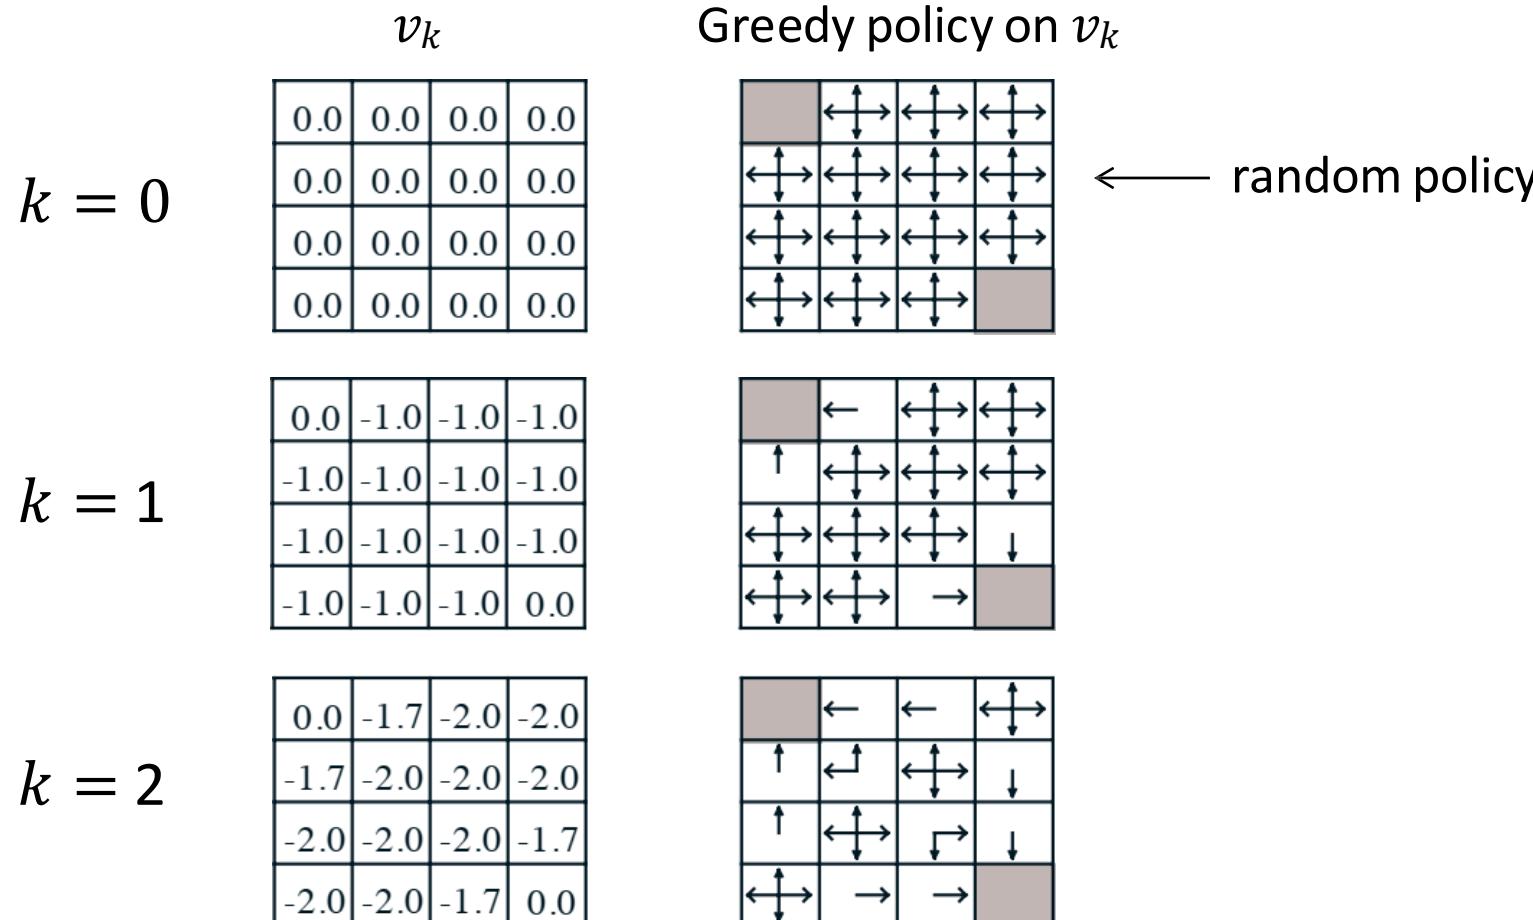
$$\pi(n| \cdot) = \pi(s| \cdot) = \pi(e| \cdot) = \pi(w| \cdot) = 0.25$$

r=1 on all transitions

	1	2	3
4	5	6	7
8	9	10	11
12	13	14	



Iterative Policy Evaluation on Small Gridworld (I)



Iterative Policy Evaluation on Small Gridworld (I)

$k = 3$

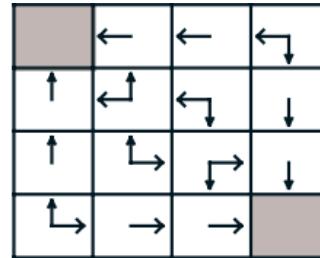
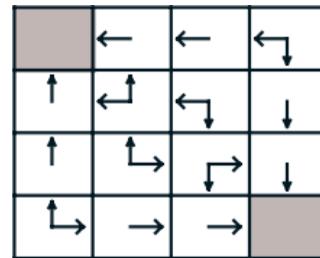
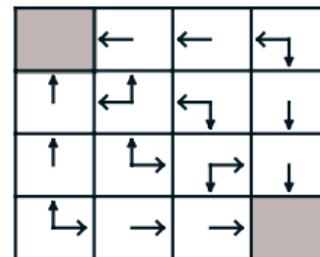
0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

$k = 10$

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0

$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0



optimal policy

Policy Iteration

How to Improve a Policy

- ✓ Given policy π
- ✓ Evaluate the policy π

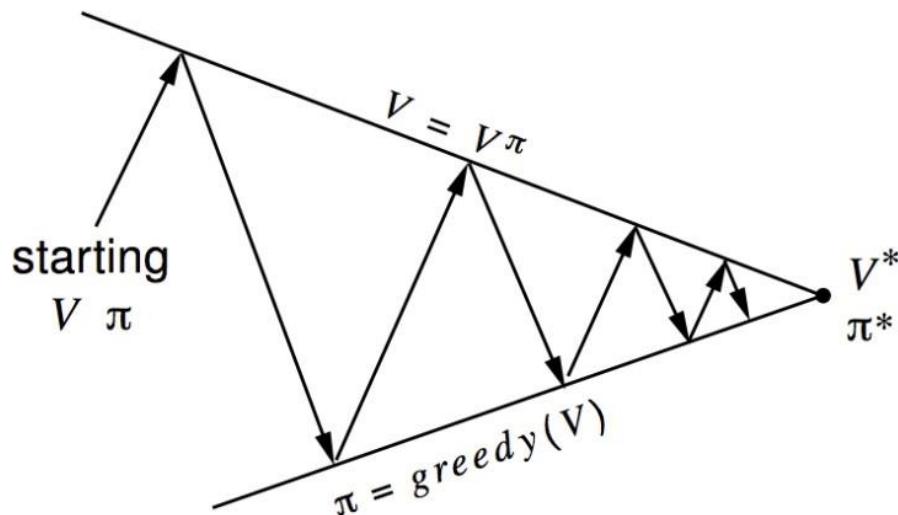
$$v_\pi(s) = \mathbb{E} [R_{t+1} + \gamma R_{t+2} + \dots | S_t = s]$$

- ✓ Improve the policy by acting greedily with respect to v_π

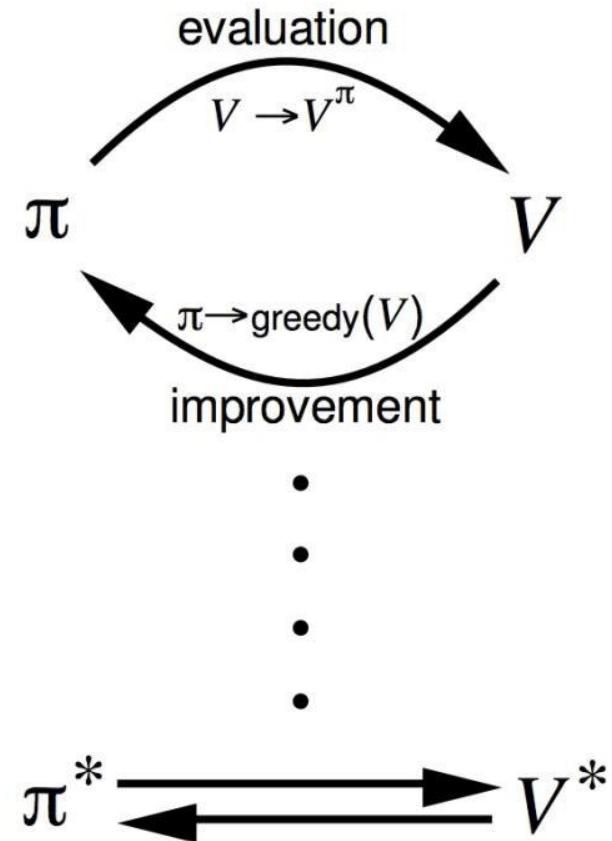
$$\pi' = \text{greedy}(\pi)$$

- ✓ In Small Gridworld improved policy was optimal, $\pi' = \pi_*$
- ✓ In general, need more iterations of improvement / evaluation
- ✓ But this process of policy iteration always converges to π_*

Policy Iteration



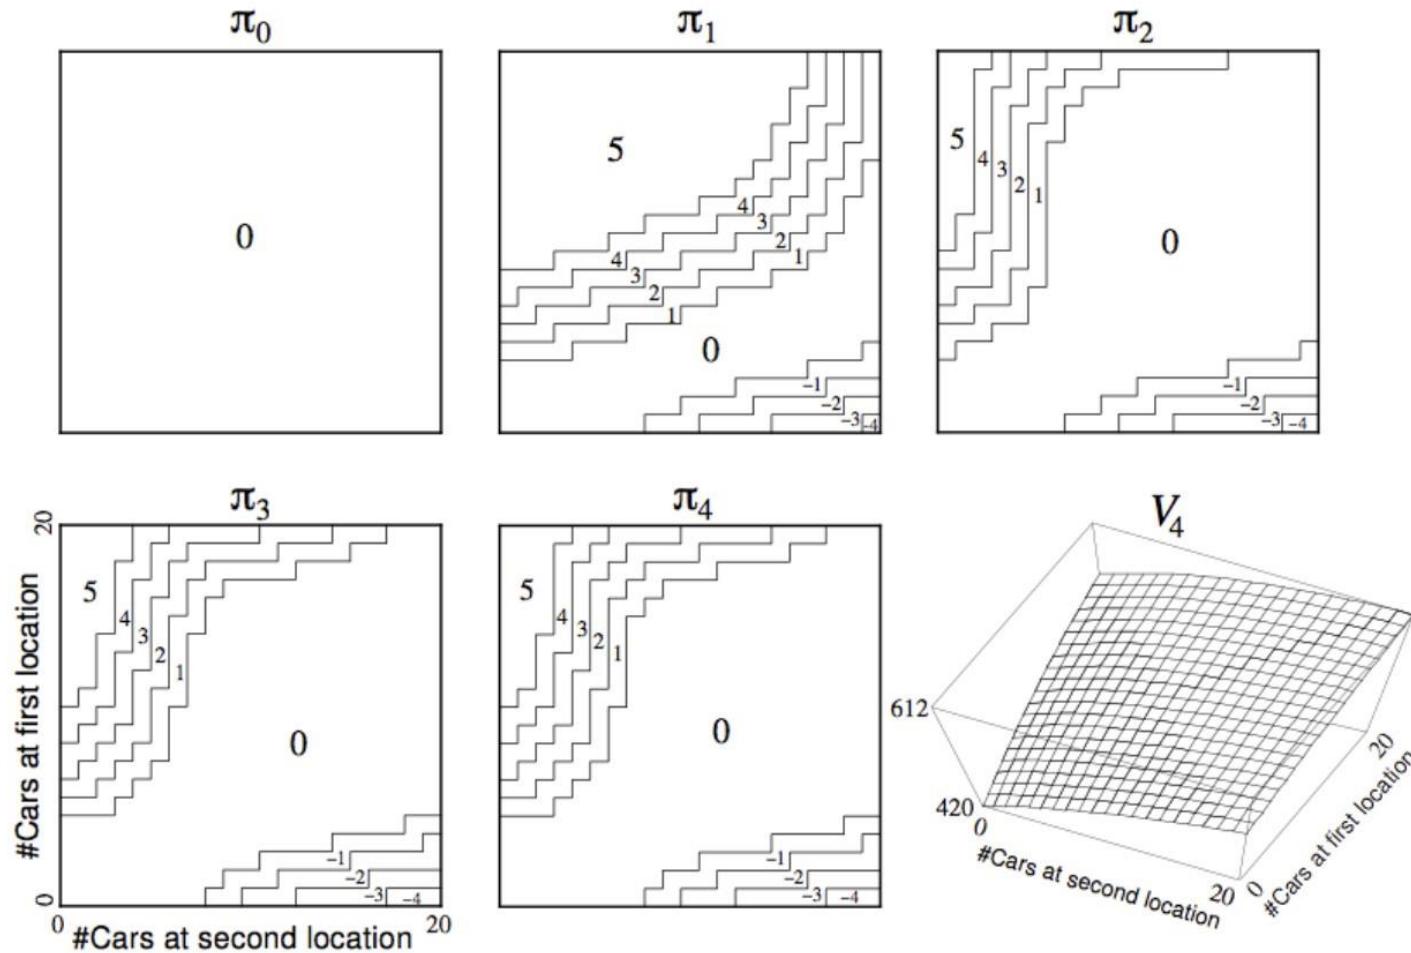
- ✓ **Policy evaluation** - Estimate v_π
- ✓ Iterative policy evaluation
- ✓ **Policy improvement** - Generate $\pi' \geq \pi$
- ✓ Greedy policy improvement





Jack's Car Rental

- ✓ States - Two locations, maximum of 20 cars at each
- ✓ Actions - Move up to 5 cars between locations overnight
- ✓ Reward - \$10 for each car rented (must be available)
- ✓ Transitions - Cars returned and requested randomly
 - ✓ Poisson distribution, n returns/requests $\sim \frac{\lambda^n}{n!} e^{-\lambda}$
 - ✓ 1st location: average requests = 3, average returns = 3
 - ✓ 2nd location: average requests = 4, average returns = 2



Policy Iteration in Jack's Car Rental

Policy Improvement (I)

Consider a deterministic policy $a = \pi(s)$

We can improve the policy by **acting greedily**

$$\pi'(s) = \arg \max_{a \in \mathcal{A}} q_\pi(s, a)$$

This **improves the value from any state s** over one step

$$q_\pi(s, \pi'(s)) = \max_{a \in \mathcal{A}} q_\pi(s, a) \geq q_\pi(s, \pi(s)) = v_\pi(s)$$

Therefore improving the value function $v_{\pi'}(s) \geq v_\pi(s)$

Policy Improvement (II)

If improvement stops

$$q_{\pi}(s, \pi'(s)) = \max_{a \in \mathcal{A}} q_{\pi}(s, a) = q_{\pi}(s, \pi(s)) = v_{\pi}(s)$$

We satisfy Bellman optimality

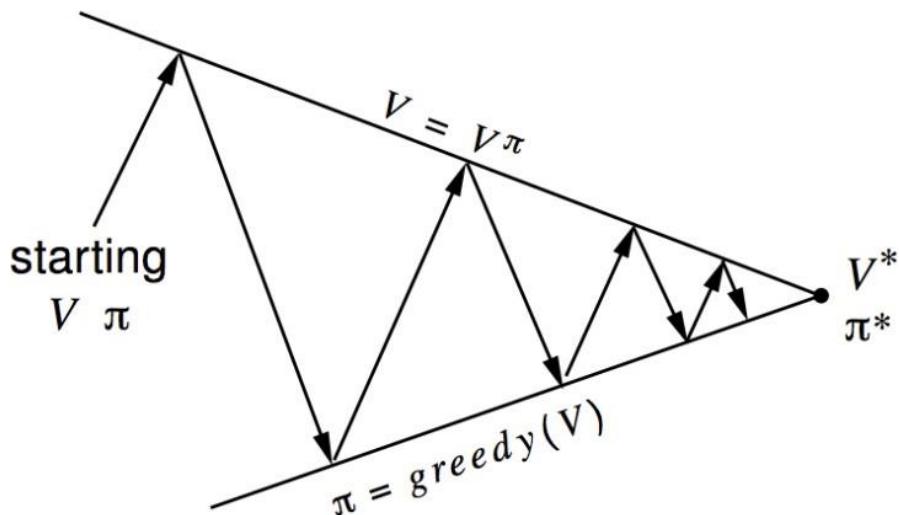
$$v_{\pi}(s) = \max_{a \in \mathcal{A}} q_{\pi}(s, a)$$

Therefore $v_{\pi}(s) = v_*(s), \forall s \in \mathcal{S}$, and π is an optimal policy

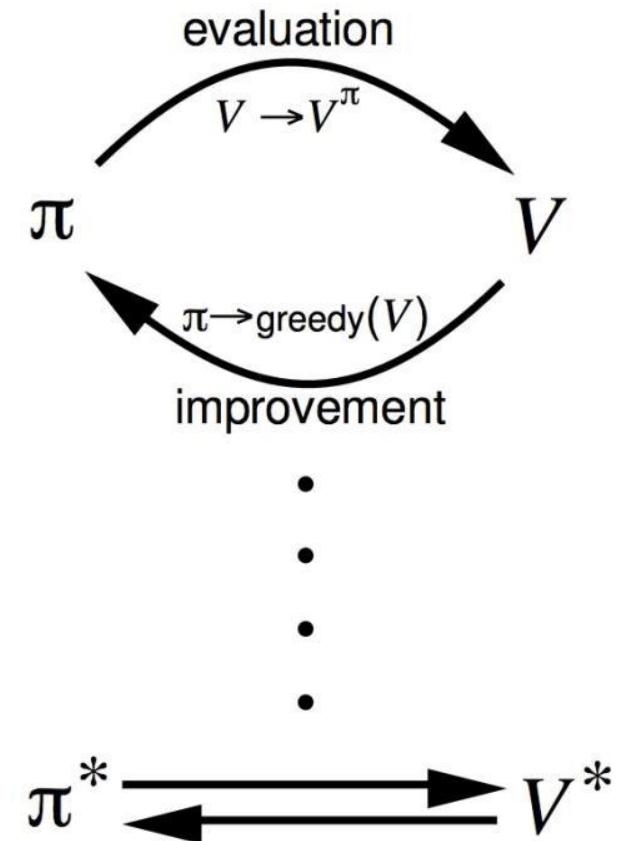
Modified Policy Improvement

- ✓ Does policy evaluation need to converge to v_{π^*} ?
 - ✓ Introduce a **stopping condition**, e.g. ϵ -convergence of value function
 - ✓ **Stop after k iterations** of iterative policy evaluation, e.g. k=3 was sufficient in small gridworld
-
- ✓ Why update policy every iteration?
 - ✓ Stop after k = 1
 - ✓ This is equivalent to value iteration (coming up)

Generalized Policy Iteration



- ✓ **Policy evaluation** - Estimate v_π
- ✓ **Any policy evaluation**
- ✓ **Policy improvement** - Generate $\pi' \geq \pi$
- ✓ **Any policy improvement algorithm**



Value Iteration

Optimality Principle

Any optimal policy can be subdivided into two components

- ✓ An optimal first action a^*
- ✓ Followed by an optimal policy from successor state s'

Theorem (Principle of Optimality)

A policy $\pi(a|s)$ achieves the optimal value from state s' (i.e. $v_\pi(s) = v_*(s)$) if and only if for any state s' reachable from s

- π achieves the optimal value from state s' , $v_\pi(s') = v_*(s')$

Deterministic Value Iteration

- ✓ If we know the solution to subproblems $v_*(s')$
- ✓ Then solution $v_*(s)$ can be found by one-step lookahead

$$v_*(s) \leftarrow \max_{a \in \mathcal{A}} \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a v_*(s')$$

- ✓ Value iteration applies these updates iteratively
- ✓ Intuition: start with final rewards and work backwards
- ✓ Still works with loopy, stochastic MDPs

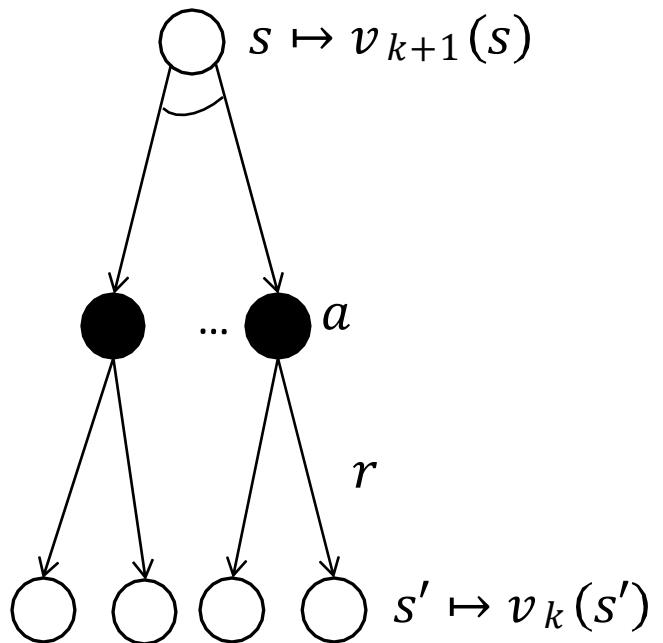
Value Iteration

- ✓ Problem: find optimal policy π
- ✓ Solution: iterative application of Bellman optimality backup

$$v_1 \rightarrow v_2 \rightarrow \cdots \rightarrow v_\pi$$

- ✓ Using synchronous backups
 - i. At each iteration $k + 1$
 - ii. For all states $s \in \mathcal{S}$
 - iii. Update $v_{k+1}(s)$ from $v_k(s')$
- ✓ Unlike policy iteration, there is no explicit policy
- ✓ Intermediate value functions may not correspond to any policy

Value Iteration - Formally



$$v_{k+1}(s) = \max_{a \in \mathcal{A}} \left(\mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a v_k(s') \right)$$

$$v_{k+1} = \max_{a \in \mathcal{A}} (\mathcal{R}^a + \gamma \mathbf{P}^a v_k)$$

The algorithm:

Initialize $V(s)$ arbitrarily, for all $s \in \mathcal{S}$

Initialize θ to a small positive value

Loop:

$$\Delta \leftarrow 0$$

Loop for each $s \in \mathcal{S}$:

$$v \leftarrow V(s)$$

$$V(s) \leftarrow \max_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$$

$$\Delta \leftarrow \max(\Delta, |v - V(s)|)$$

Until $\Delta < \theta$

Output a deterministic policy, $\pi \approx \pi^*$, such that

$$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$$

DP Example

https://cs.stanford.edu/people/karpathy/reinforcejs/gridworld_dp.html

Synchronous Dynamic Programming

Wrap-up

Problem	Bellman Equation	Algorithm
Prediction	Bellman Expectation Equation	Iterative Policy Evaluation
Control	Bellman Expectation Equation + Greedy Policy Improvement	Policy Iteration
Control	Bellman Optimality Equation	Value Iteration

- ✓ Algorithms are based on **state-value function** $v_\pi(s)$ or $v_*(s)$
- ✓ Complexity is $O(mn^2)$ per iteration ($m = |\mathcal{A}|$ and $n = |\mathcal{S}|$)
- ✓ Could also apply to **action-value function** $q_\pi(s, a)$ or $q_*(s, a)$
- ✓ Complexity is $O(m^2n^2)$ per iteration

Take (stay) home messages

- ✓ **Dynamic Programming** - Method for solving complex problems by breaking them down into subproblems
 - ✓ Use recursive formulation founded in [return nested definition](#)
- ✓ **Policy iteration** - Re-define the policy at each step and compute the value according to this new policy until the policy converges
- ✓ **Value iteration** - Computes the optimal state value function by iteratively improving the estimate of $V(s)$
- ✓ **Policy vs Value** iteration
 - ✓ Policy can converge quicker (agent is interested in optimal policy)
 - ✓ Value iteration is computationally cheaper (per iteration)

Next Lecture

Model-Free Prediction

- ✓ Estimate the **value function** of an unknown MDP
- ✓ Monte-Carlo approaches
- ✓ Temporal-Difference learning
- ✓ $\text{TD}(\lambda)$

Before start...

Primo appello sessione estiva: 16/06

Secondo appello sessione estiva: 01/07 (NEW!)

Terzo appello sessione estiva: 21/07

- Domani Tezione speciale AIRC campus: uso degli strumenti computazionali per lo studio del cancro.
- — — — —



Model Free Prediction

Introduction

Outline

- ✓ Introduction
- ✓ Monte-Carlo approaches
- ✓ Temporal-Difference (TD) learning
- ✓ TD(λ)

Model-Free Reinforcement Learning

- ✓ So far: solve a **known MDP** (states, transition, rewards, actions)
- ✓ Model free
 - ✓ No environment model
 - ✓ **No knowledge of MDP** transition/rewards
- ✓ **Model-free prediction** - Estimate the value function of an unknown MDP
- ✓ **Model-free control** - Optimise the value function of an unknown MDP

Monte-Carlo

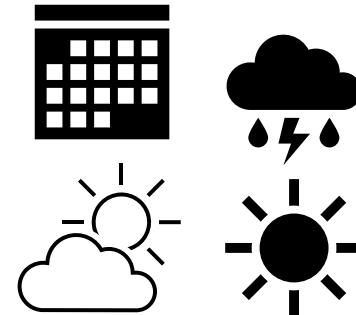
Monte-Carlo (MC) Reinforcement Learning

- ✓ MC methods learn directly **from episodes of experience**
- ✓ MC is **model-free**: no knowledge of MDP transitions/rewards
- ✓ MC learns from **complete episodes**: no bootstrapping

Bootstrapping in RL

Model that, on Monday, predicts the meteo for Sunday.

- **No bootstrap:** Takes only info from Monday and Sunday
- **Bootstrap:** Uses information from the week meteo
(induces bias)



Monte-Carlo (MC) Reinforcement Learning

- ✓ MC methods learn directly **from episodes of experience**
- ✓ MC is **model-free**: no knowledge of MDP transitions/rewards
- ✓ MC learns from **complete episodes**: no bootstrapping(*)
- ✓ MC uses the simplest possible idea: value = mean return across episodes
- ✓ Caveat: can only apply MC to episodic MDPs
 - ✓ All **episodes must terminate**

Monte-Carlo Policy Evaluation

- ✓ Goal: learn v_π from episodes of experience under policy π

$$S_1, A_1, R_2, \dots, R_k \sim \pi$$

- ✓ Recall that return is the total discounted reward

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$$

- ✓ Recall that value function is the expected return

$$v_\pi(s) = \mathbb{E}[G_t | S_t = s]$$

- ✓ Monte-Carlo policy evaluation uses empirical mean return instead of expected return

First-Visit Monte-Carlo Policy Evaluation

- ✓ To evaluate state s
- ✓ The first time-step t that state s is visited in an episode
 - I. Increment counter $N(s) \leftarrow N(s) + 1$
 - II. Increment total return $T(s) \leftarrow T(s) + G_t$
 - III. Value is estimated by mean return $V(s) = T(s)/N(s)$
- ✓ By law of large numbers

$$V(s) \rightarrow v_\pi(s) \text{ as } N(s) \rightarrow \infty$$

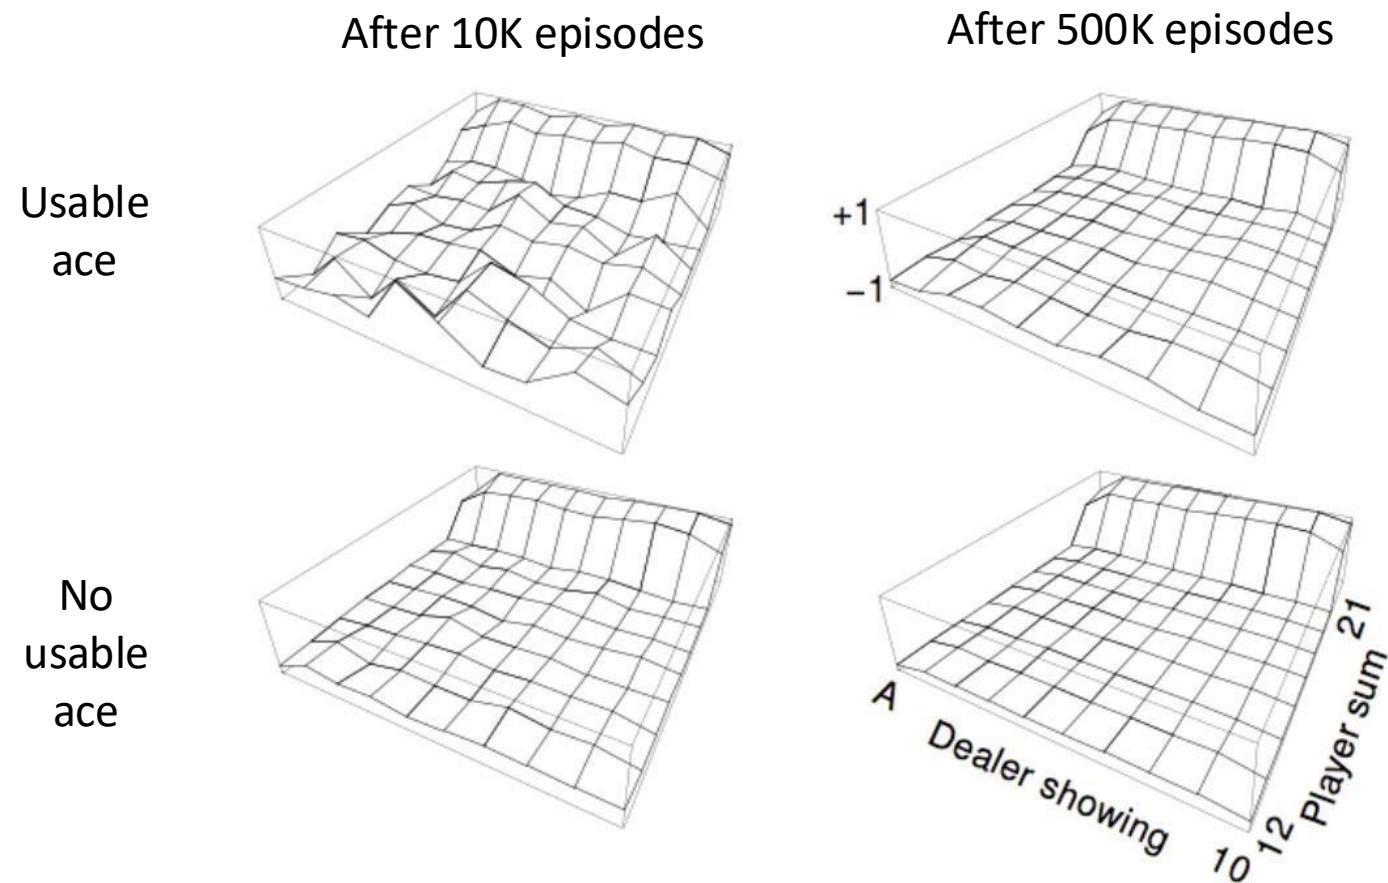
Every-Visit Monte-Carlo Policy Evaluation

- ✓ To evaluate state s
- ✓ Every time-step t that state s is visited in an episode
 - I. Increment counter $N(s) \leftarrow N(s) + 1$
 - II. Increment total return $T(s) \leftarrow T(s) + G_t$
 - III. Value is estimated by mean return $V(s) = T(s)/N(s)$



Blackjack Example

Blackjack Value Function after MC Learning



Policy: **stick** if sum of cards ≥ 20 , otherwise **twist**

Incremental Mean

The mean μ_1, μ_2, \dots of a sequence x_1, x_2, \dots can be computed incrementally

$$\mu_k = \frac{1}{k} \bar{\mu} \sum_{j=1}^k x_j = \frac{1}{k} \left(x_k + \sum_{j=1}^{k-1} x_j \right)$$

$$\mu_k = \frac{1}{k} (x_k + (k-1)\mu_{k-1}) = \mu_{k-1} + \frac{1}{k} (x_k - \mu_{k-1})$$

Incremental Mean MC Update

✓ Update $V(s)$ incrementally after episode $S_1, A_1, R_2, \dots, R_T$

✓ For each state S_t with return G_t

I. Increment counter $N(s) \leftarrow N(s) + 1$

II. Update value function (with incremental mean)

$$V(S_t) \leftarrow V(S_t) + \frac{1}{N(S_t)} (G_t - V(S_t))$$

✓ In **non-stationary problems** track a running mean (forget old episodes)

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$

Temporal-Difference Learning

Temporal-Difference (TD) Learning

- ✓ TD methods learn directly from episodes of experience
- ✓ TD is model-free: no knowledge of MDP transitions / rewards
- ✓ TD **learns from incomplete episodes**, by bootstrapping
- ✓ TD **updates a guess towards a guess**

MC Vs TD Learning

✓ Goal: learn v_π from episodes of experience under policy π

✓ Incremental every-visit MC

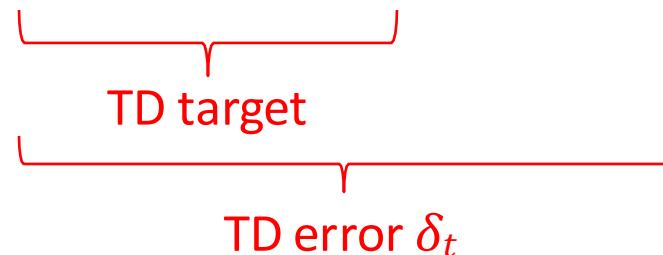
✓ Update value $V(S_t)$ toward actual return G_t

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$

✓ Simplest temporal-difference learning algorithm (TD(0))

✓ Update value $V(S_t)$ toward estimated return $R_t + \gamma V(S_{t+1})$

$$V(S_t) \leftarrow V(S_t) + \alpha(R_t + \gamma V(S_{t+1}) - V(S_t))$$



TD(0) Learning algorithm

Algorithm 1 Tabular TD(0) for estimating v_π

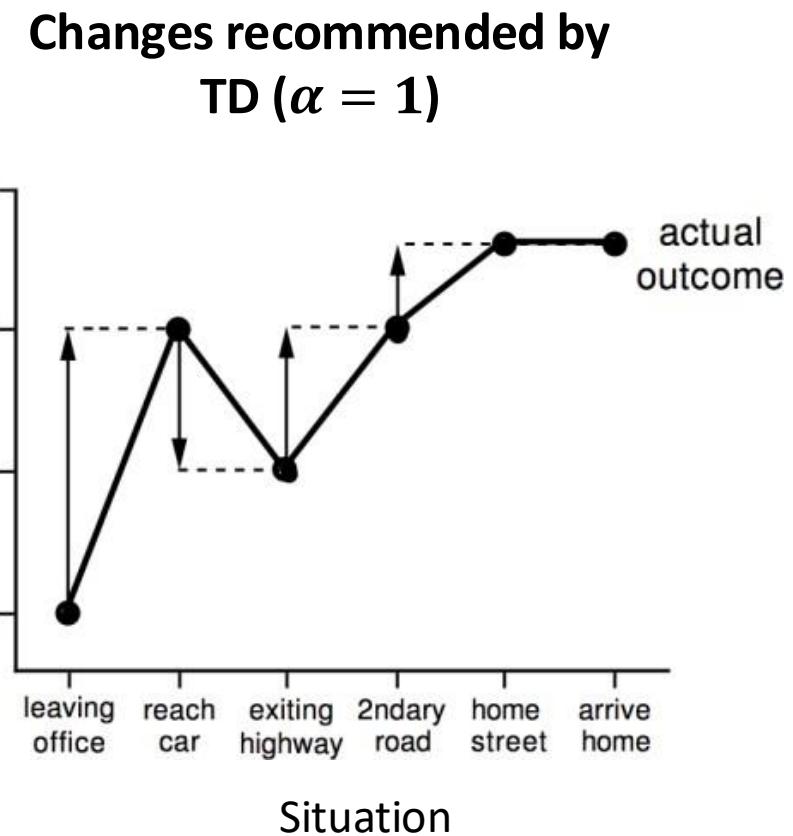
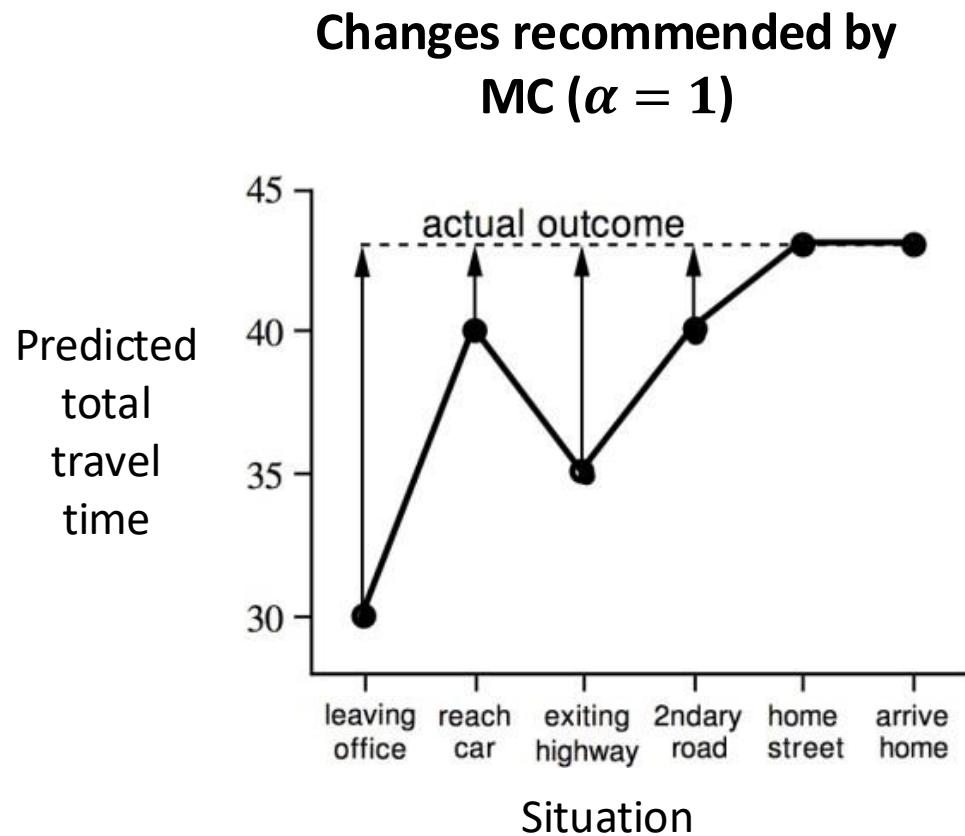
Input: Policy π to be evaluated **Parameters:** Learning rate $\alpha \in (0, 1]$

```
1: for each episode: do
2:   Initialize  $S$ 
3:   while  $S$  is not terminal: do
4:     Take action  $A$  given by  $\pi(a|S)$ 
5:     Observe  $R, S'$ 
6:     Update  $V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$ 
7:      $S \leftarrow S'$ 
8:   end while
9: end for
```

Driving Home Example

State	Elapsed Time (minutes)	Predicted Time to Go	Predicted Total Time
leaving office	0	30	30
reach car, raining	5	35	40
exit highway	20	15	35
behind truck	30	10	40
home street	40	3	43
arrive home	43	0	43

Driving Home Example – MC vs TD



Advantages and Disadvantages of MC vs. TD (I)

- ✓ TD can learn **before** knowing the final outcome
 - ✓ TD can learn online after every step
 - ✓ MC must wait until end of episode before return is known
- ✓ TD can learn **without** the final outcome
 - ✓ TD can learn from incomplete sequences
 - ✓ MC can only learn from complete sequences
 - ✓ TD works in continuing (non-terminating) environments
 - ✓ MC only works for episodic (terminating) environments

Bias-Variance Tradeoff

- ✓ Return $G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$ is **unbiased estimate** of $v_\pi(S_t)$
- ✓ True TD target $R_{t+1} + \gamma v_\pi(S_{t+1})$ is **unbiased estimate** of $v_\pi(S_t)$
- ✓ TD target $R_{t+1} + \gamma V(S_{t+1})$ is **biased estimate** of $v_\pi(S_t)$
- ✓ TD target is much lower variance than the return:
 - ✓ Return depends on many random actions, transitions, rewards
 - ✓ TD target depends on one random action, transition, reward

Advantages and Disadvantages of MC vs. TD (II)

- ✓ MC has **high variance, zero bias**
 - ✓ Good convergence properties (even with function approximation)
 - ✓ Not very sensitive to initial value
 - ✓ Very simple to understand and use

- ✓ TD has **low variance, some bias**
 - ✓ Usually more efficient than MC
 - ✓ TD(0) converges to $v_\pi(s)$ (but not always with function approximation)
 - ✓ More sensitive to initial value

Batch MC and TD

- ✓ MC and TD **converge**: $V(s) \rightarrow v_\pi(s)$ as experience $\rightarrow \infty$
- ✓ But what about **batch solution for finite experience**?

$$\begin{aligned} s_1^1, a_1^1, r_2^1, \dots, s_{T_1}^1 \\ \vdots \\ s_1^K, a_1^K, r_2^K, \dots, s_{T_K}^K \end{aligned}$$

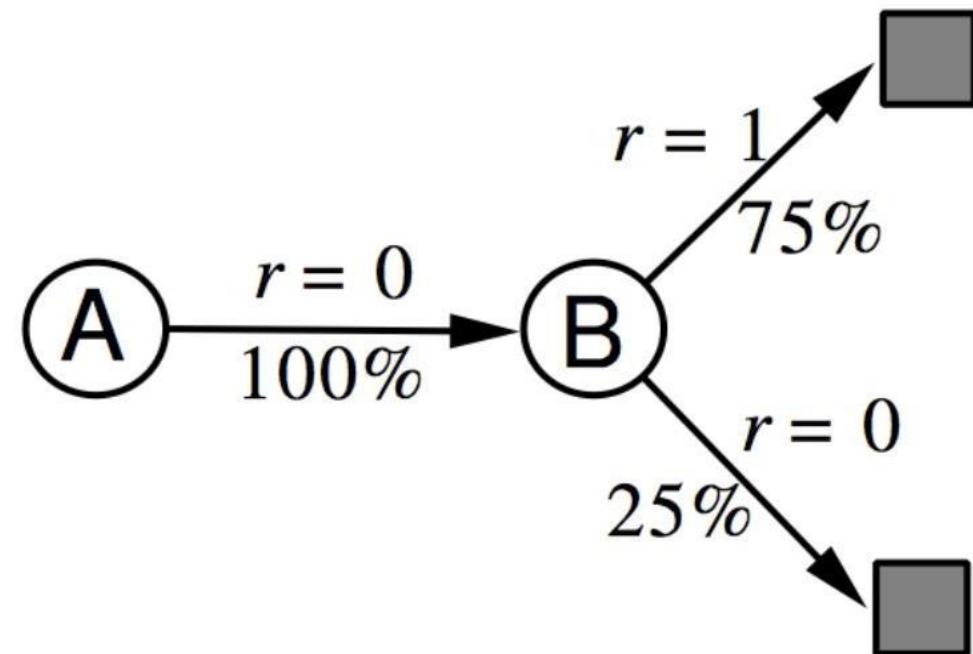
- ✓ e.g. repeatedly sample episode $k \in [1, K]$
- ✓ Apply MC or TD(0) to episode k

A Simple Example

✓ Two states A; B; no discounting; 8 episodes of experience

1. A, 0, B, 0
2. B, 1
3. B, 1
4. B, 1
5. B, 1
6. B, 1
7. B, 1
8. B, 0

✓ What is $V(A)$; $V(B)$?



Certainty Equivariance

- ✓ MC converges to solution with **minimum mean-squared error**
- ✓ Best fit to the observed returns

$$\sum_{k=1}^K \sum_{t=1}^{T_k} (G_t^k - V(s_t^k))^2$$

- ✓ TD(0) converges to solution of **maximum likelihood Markov model**
- ✓ Solution to the MDP $\langle \mathcal{S}, \mathcal{A}, \mathbf{P}, \mathcal{R}, \gamma \rangle$ that best fits the data

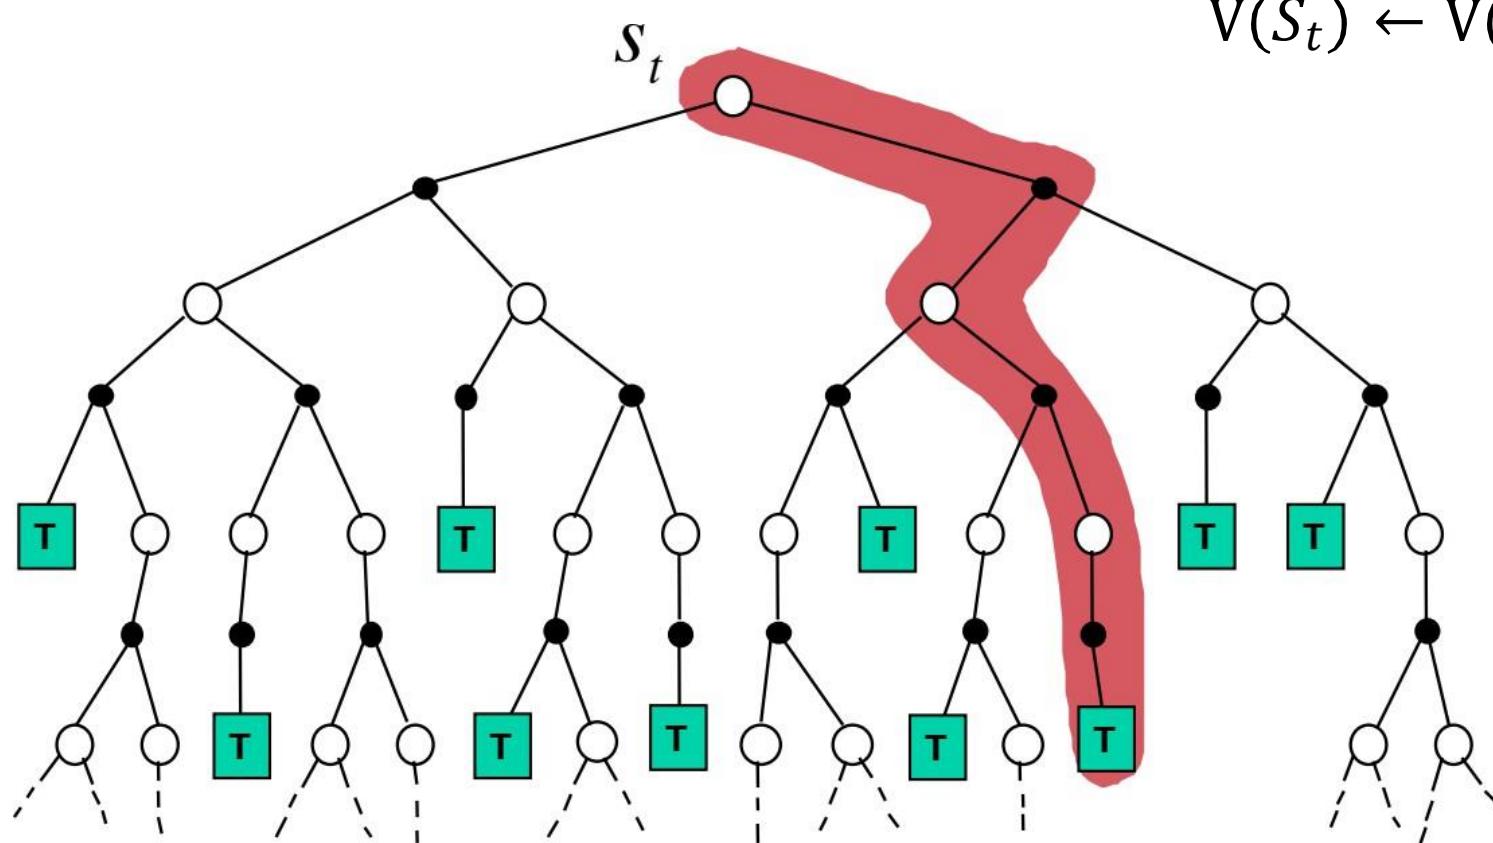
$$\hat{P}_{ss'}^a = \frac{1}{N(s, a)} \sum_{k=1}^K \sum_{t=1}^{T_k} \mathbf{1}(s_t^k, a_t^k, s_{t+1}^k; s, a, s')$$
$$\hat{\mathcal{R}}_s^a = \frac{1}{N(s, a)} \sum_{k=1}^K \sum_{t=1}^{T_k} \mathbf{1}(s_t^k, a_t^k; s, a) r_t^k$$

Advantages and Disadvantages of MC vs. TD (III)

- ✓ TD exploits **Markov property**
- ✓ Usually more efficient in Markov environments
- ✓ MC does **not exploit Markov** property
- ✓ Usually more effective in non-Markov environments

Unified View

MC Update

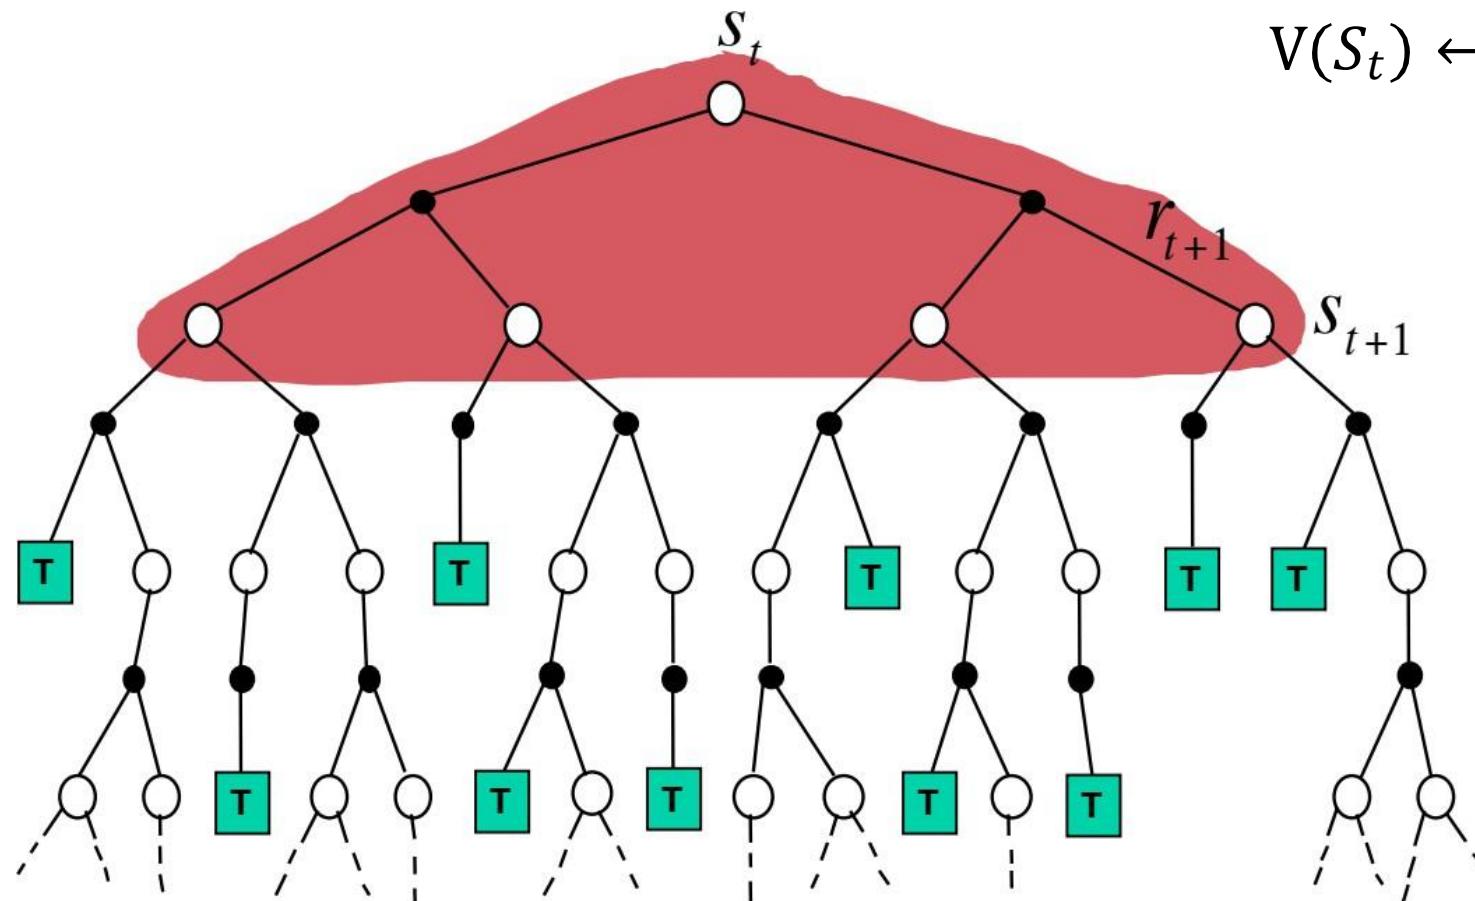


$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$

TD Update

$$V(S_t) \leftarrow V(S_t) + \alpha(R_t + \gamma V(S_{t+1}) - V(S_t))$$

Dynamic Programming



$$V(S_t) \leftarrow \mathbb{E}[R_{t+1} + \gamma V(S_{t+1})]$$

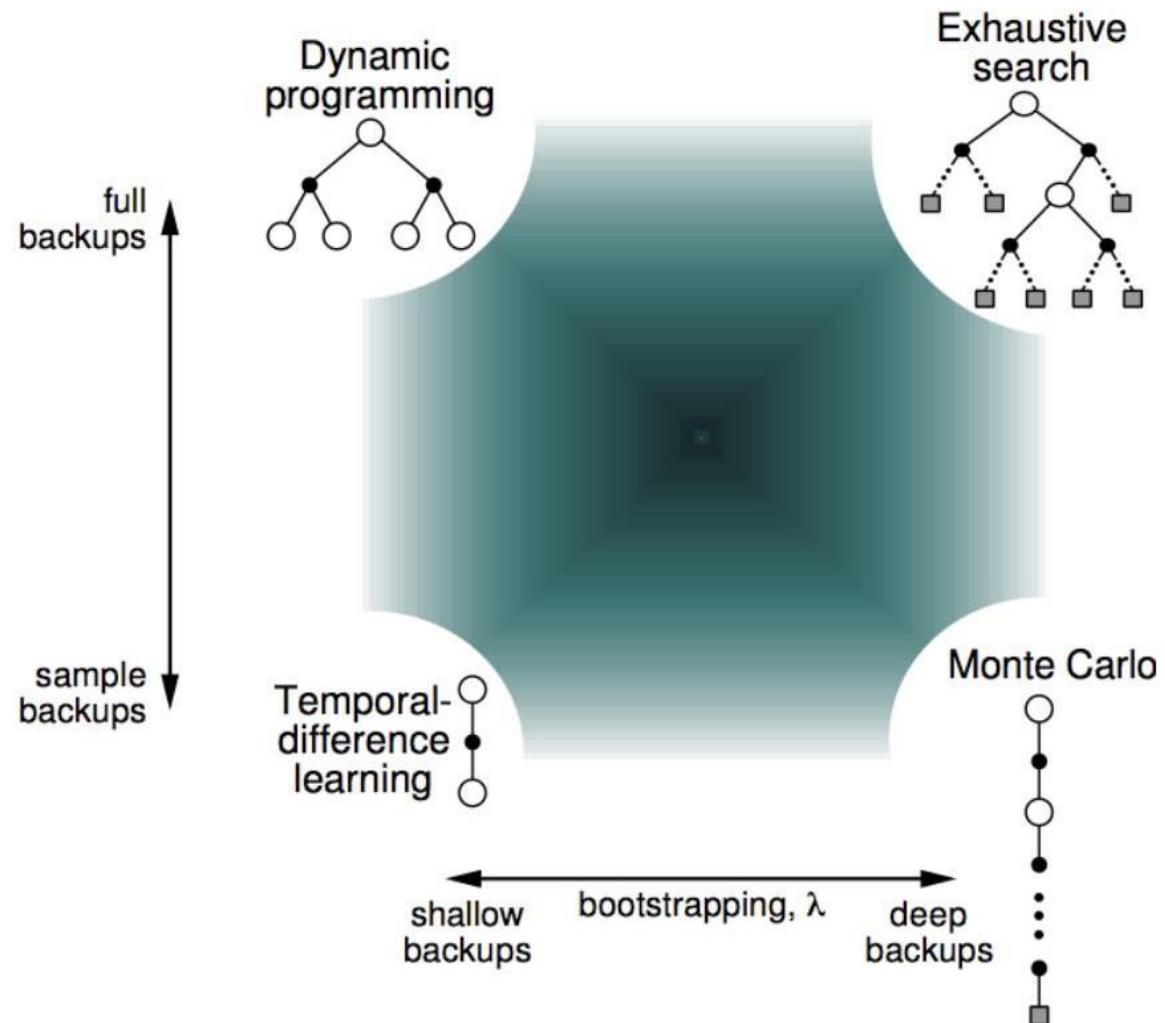
Bootstrapping and Sampling

✓ Bootstrapping - Update involves an estimate

- ✓ MC does not bootstrap
- ✓ DP bootstraps
- ✓ TD bootstraps

✓ Sampling - Update samples an expectation

- ✓ MC samples
- ✓ DP does not sample
- ✓ TD samples

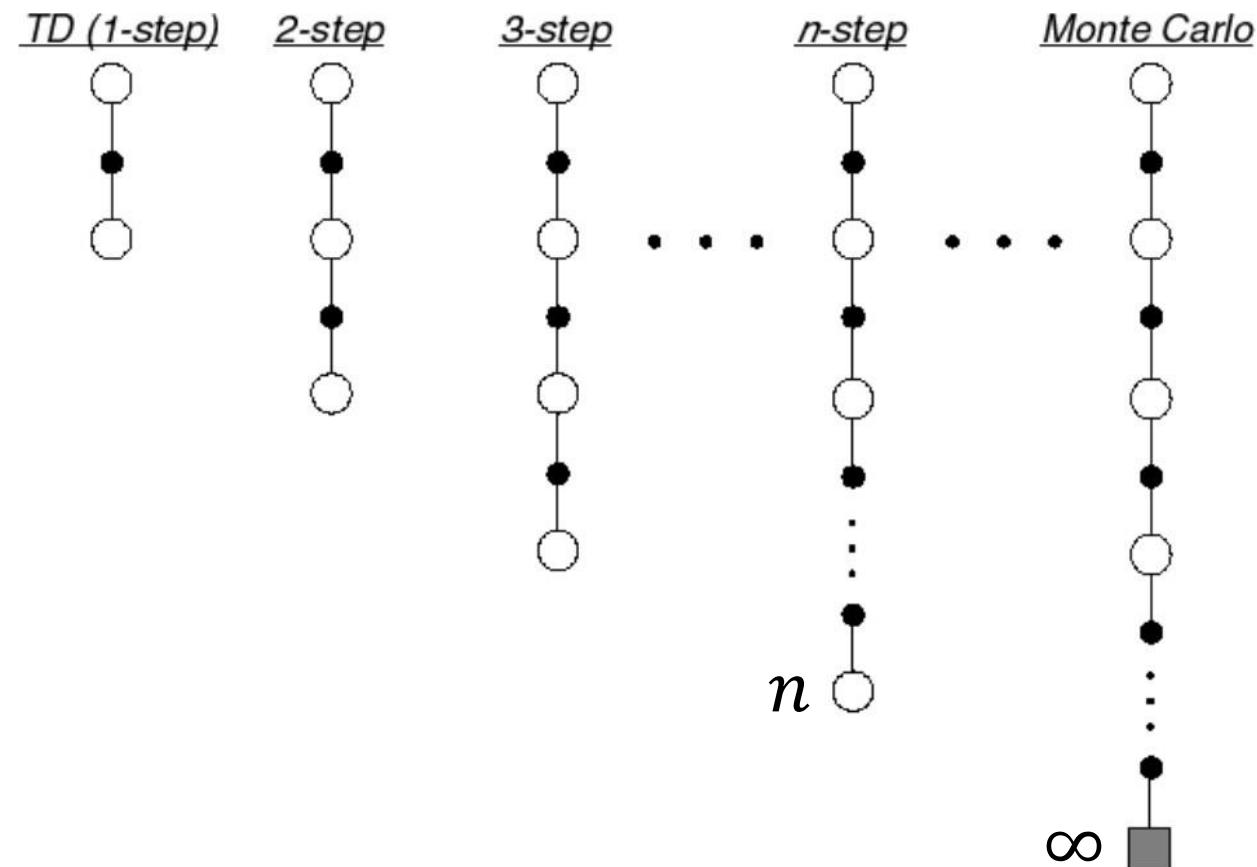


Unified View of RL

Generalizing TD

n -step Prediction

Have TD look and target n steps in the future



n-step Return

- ✓ Consider the following *n* -step returns for $n = 1, 2, \dots, \infty$

$$\begin{aligned} n = 1 \quad (\text{TD}) \quad G_t^{(1)} &= R_{t+1} + \gamma V(S_{t+1}) \\ n = 2 \quad &\quad G_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma V(S_{t+2}) \\ &\quad \dots \\ n = \infty \quad (\text{MC}) \quad G_t^{(\infty)} &= R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T \end{aligned}$$

- ✓ Define the *n*-step return

$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n})$$

- ✓ Learn based on the *n*-step difference

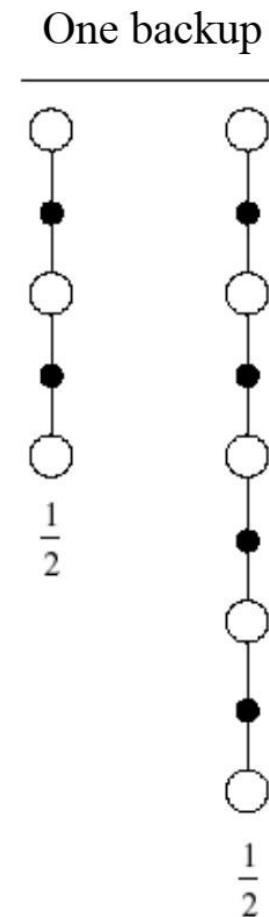
$$V(S_t) \leftarrow V(S_t) + \alpha \left(G_t^{(n)} - V(S_t) \right)$$

Averaging n -step Returns

- ✓ We can average n -step returns over different n
 - ✓ E.g.: Average the 2-step and 4-step returns

$$\frac{1}{2} G^{(2)} + \frac{1}{4} G^{(4)}$$

- ✓ Combines information from two different time-steps
- ✓ Can we efficiently combine information from all time-steps?



λ -returns

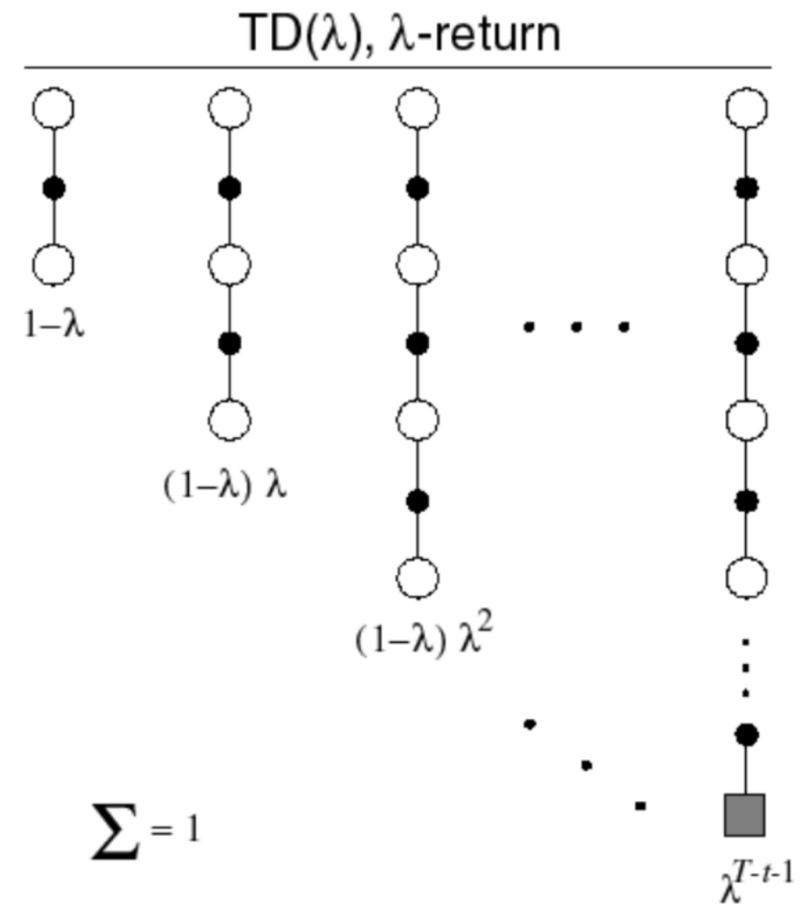
- ✓ The λ -return G_t^λ combines all n -step returns $G_t^{(n)}$

- ✓ Using weight $(1 - \lambda)\lambda^{n-1}$

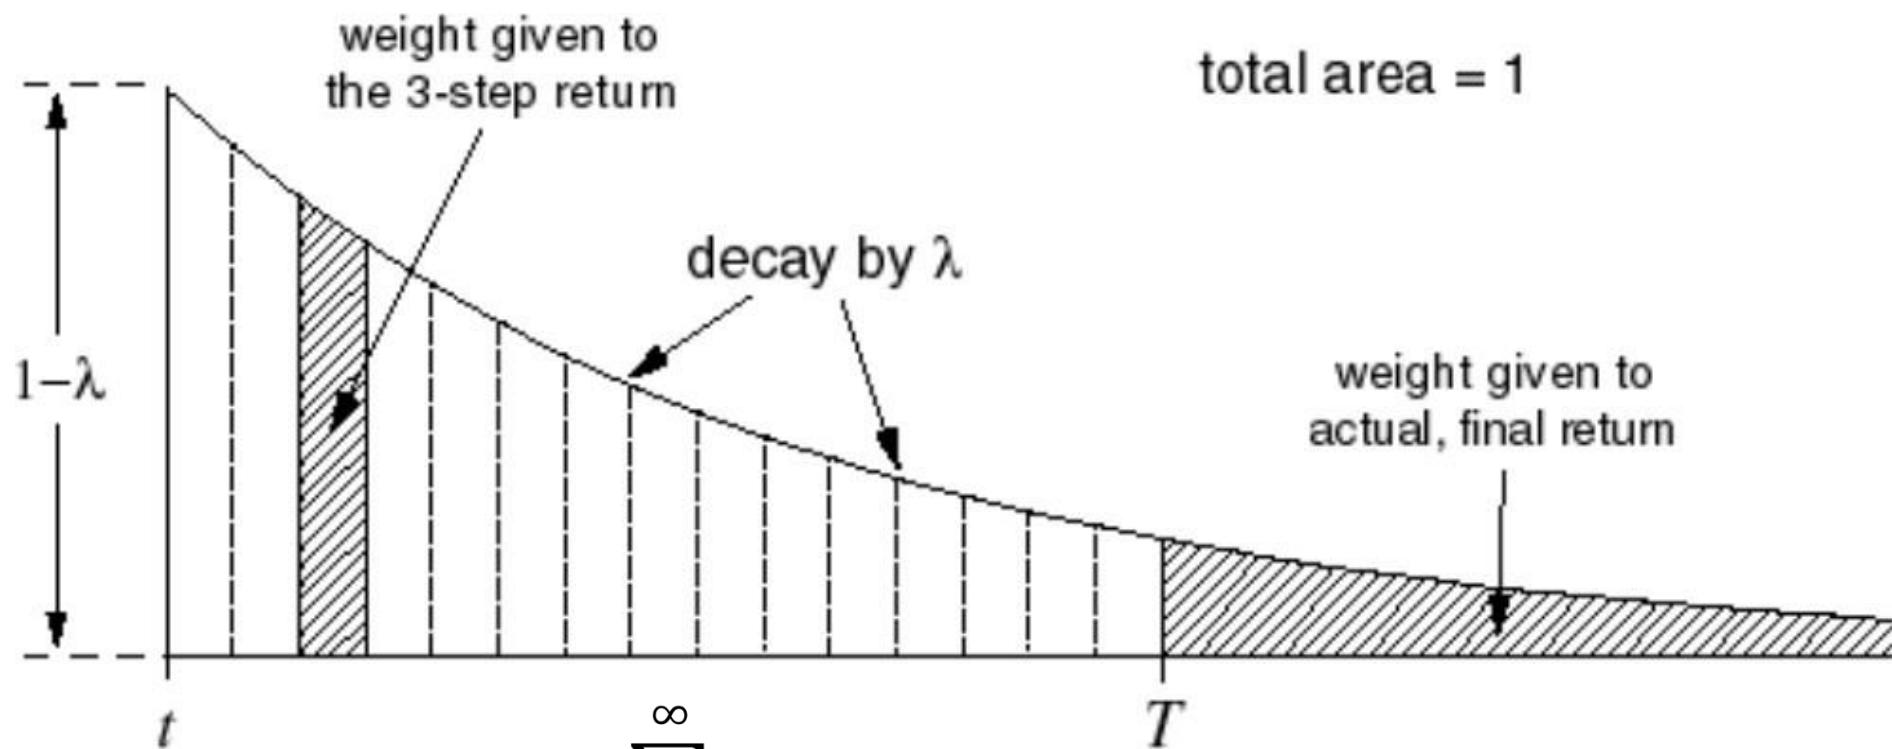
$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

- ✓ Update as appropriate (TD(λ))

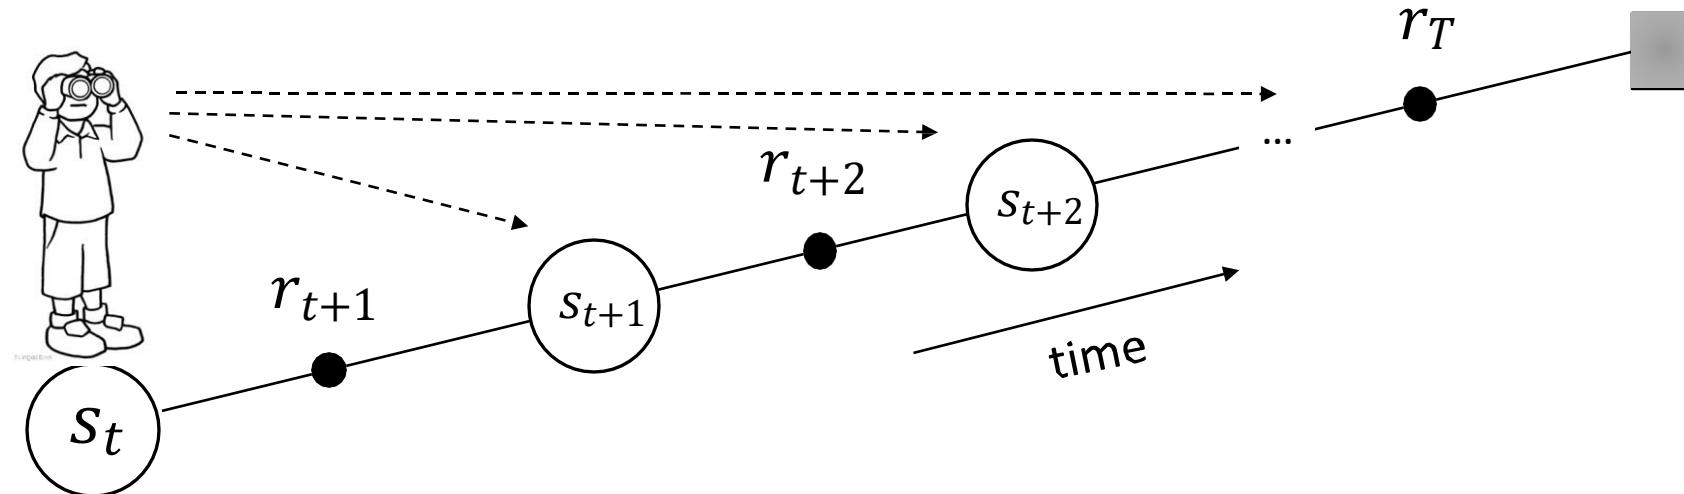
$$V(S_t) \leftarrow V(S_t) + \alpha(G_t^\lambda - V(S_t))$$



TD(λ) Weight Function



Forward View TD(λ)

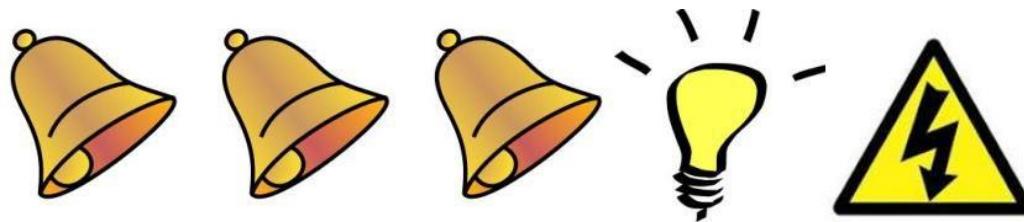


- ✓ Update value function towards the λ -return
- ✓ Forward-view looks **into the future to compute G_t^λ**
- ✓ Like MC, can only be computed from **complete episodes**

Backward View TD(λ)

- ✓ Forward view provides theory
- ✓ Backward view provides mechanism
- ✓ Update online, every step, from incomplete sequences

Eligibility Traces

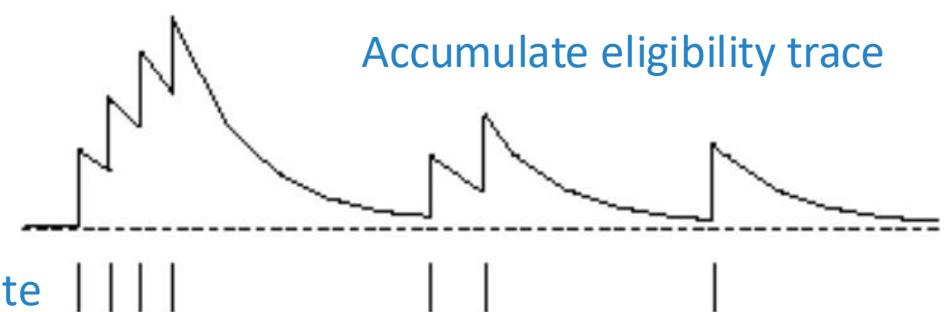


- ✓ Credit assignment problem: what caused shock?
- ✓ Frequency heuristic: assign credit to most frequent states
- ✓ Recency heuristic: assign credit to most recent states
- ✓ Eligibility traces combine both heuristics

$$E_0(s) = 0$$

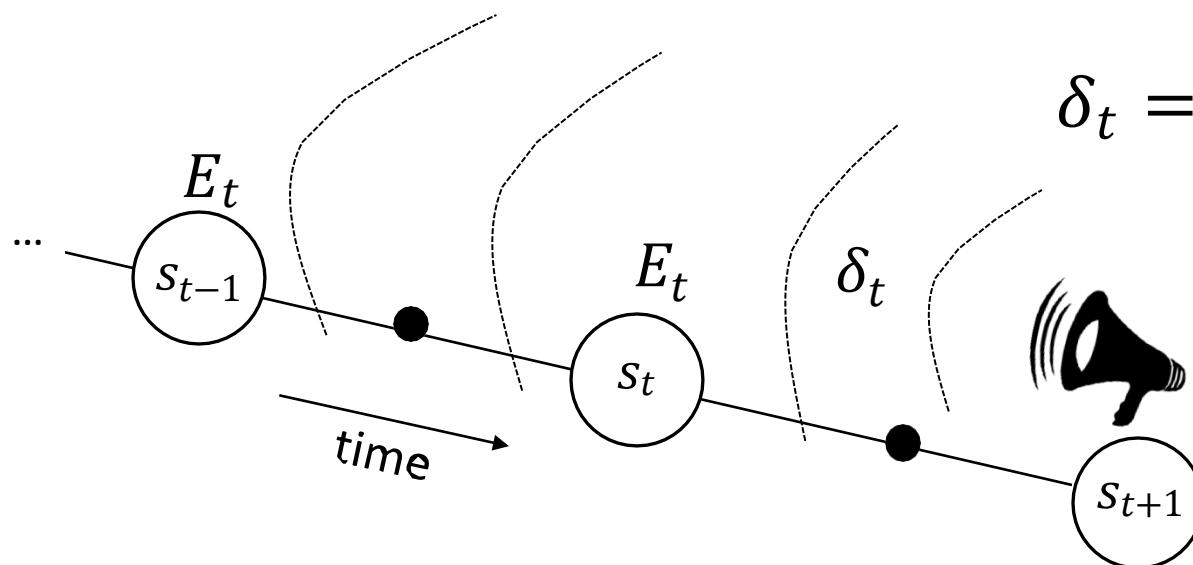
$$E_t(s) = \gamma \lambda E_{t-1}(s) + \mathbf{1}(S_t; s)$$

times of visit to state



Backward View TD(λ)

- ✓ Keep an **eligibility trace** for every state s
- ✓ Update value $V(s)$ for every state s
- ✓ In proportion to **TD-error** δ_t and **eligibility trace** $E_t(s)$



$$\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$$

$$V(s) = V(s) + \alpha \delta_t E_t(s)$$

TD(λ) and TD(0)

- ✓ When $\lambda = 0$ only current state is updated

$$E_t(s) = \mathbf{1}(S_t; s)$$

$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$

- ✓ Equivalent to TD(0) update

$$V(S_t) \leftarrow V(S_t) + \alpha \delta_t$$

TD(λ) and MC

- ✓ When $\lambda = 1$ credit is deferred until end of episode
- ✓ Consider episodic environments with offline updates
- ✓ Over the course of an episode, total update for TD(1) is the same as total update for MC

Theorem

The sum of offline updates is identical for forward-view and backward-view TD(λ)

$$\sum_{t=1}^T \alpha \delta_t E_t(s) = \sum_{t=1}^T \alpha (G_t^\lambda - V(S_t)) \mathbf{1}(S_t; s)$$

Telescoping in TD(1)

- ✓ When $\lambda = 1$ sum of TD errors telescopes into MC error
... (proof in book if interested) ...
- ✓ TD(1) is roughly equivalent to every-visit Monte-Carlo
- ✓ Error is accumulated online, step-by-step
- ✓ If value function is only updated offline at end of episode, then total update is the same as MC

Wrap-up

Take home messages

- ✓ Model-free prediction is **value function estimation of an unknown MDP**
 - ✓ Based on **sample-updates**
- ✓ **Monte Carlo** methods
 - ✓ Estimating value function by averaging sample returns
 - ✓ Only for episodic tasks (eventually terminate no matter what actions are taken)
- ✓ **TD learning**
 - ✓ Learn from existing (biased) estimates of future return (**bootstrapping**)
 - ✓ Explore the future until n-th step

Next Lecture

Model-Free Control

- ✓ Optimise the value function of an unknown MDP
- ✓ Generalised Policy Iteration
- ✓ Monte Carlo Control
- ✓ TD learning
- ✓ On-policy Vs Off-policy

Model Free Control

Outline

- ✓ Introduction
- ✓ On-policy Vs Off-Policy
- ✓ On-policy Monte-Carlo
- ✓ On-policy TD learning (SARSA)
- ✓ Off-policy TD (Q-learning)

Introduction

Today's focus

- ✓ Last lecture
 - ✓ Model-free prediction
 - ✓ Estimate the value function of an unknown MDP
-
- ✓ Today's lecture
 - ✓ Model-free control
 - ✓ Optimise the value function of an unknown MDP

Model Free Control – Where to find it

- ✓ Elevator
- ✓ Robot walking
- ✓ Vehicle Steering
- ✓ Bioreactor
- ✓ Molecule engineering
- ✓ Robocup Soccer
- ✓ Quake
- ✓ Portfolio management
- ✓ Protein Folding
- ✓ Game of Go

For most of these problems, either:

- ✓ MDP model is unknown, but experience can be sampled
- ✓ MDP model is known, but is too big to use, except by samples

Model-free control can solve these problems

On-policy & Off-policy Learning

✓ On-policy learning

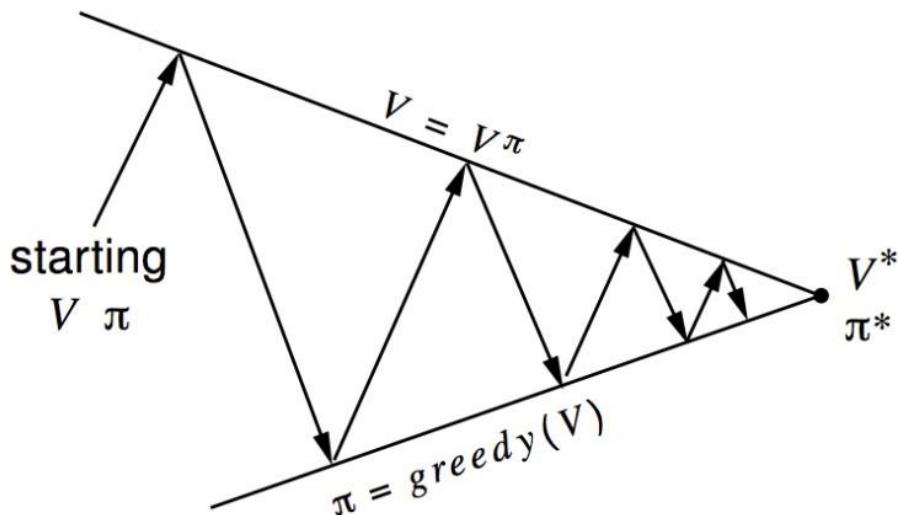
- ✓ *Learn on the job*
- ✓ Learn about policy π from experience sampled from π

✓ Off-policy learning

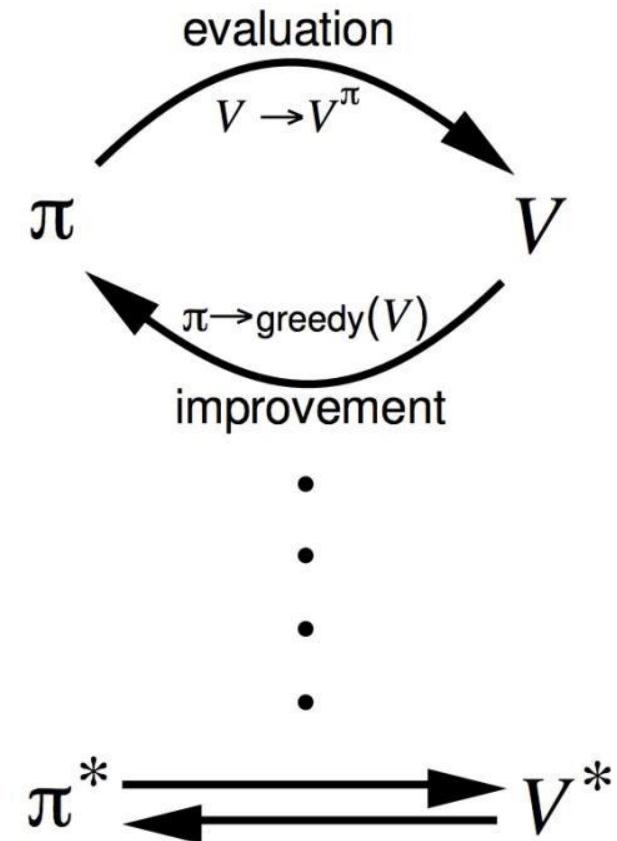
- ✓ *Look over someone's shoulder*
- ✓ Learn about policy π from experience sampled from μ

On-policy MC

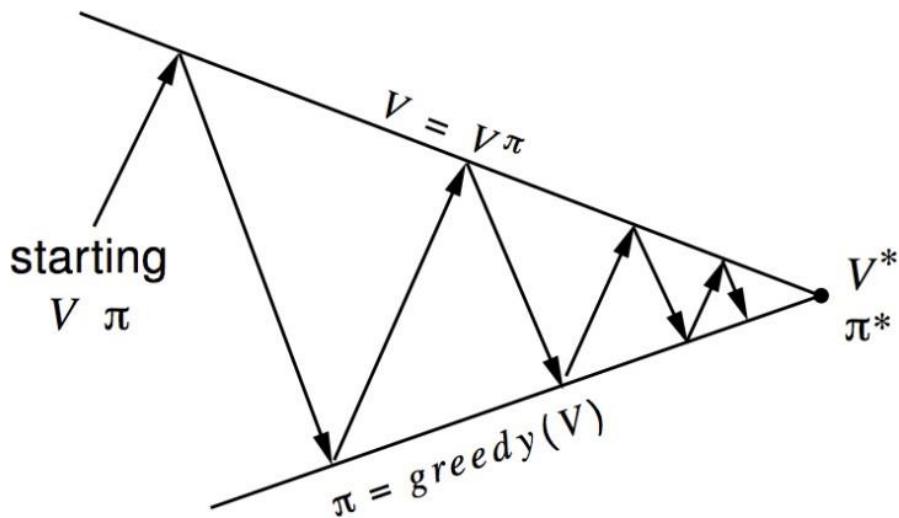
Generalized Policy Iteration (Lecture 3)



- ✓ **Policy evaluation** - Estimate v_π
- ✓ **Any policy evaluation**
- ✓ **Policy improvement** - Generate $\pi' \geq \pi$
- ✓ **Any policy improvement algorithm**



Generalized Policy Iteration with On-policy MC



- ✓ Policy evaluation - Monte-Carlo policy evaluation, $V = v_\pi$?
- ✓ Policy improvement - Generate greedy policy improvement?

Model-Free Policy Iteration Using Action-Value Function

- ✓ Greedy policy improvement over $V(s)$ requires model of MDP

$$\pi'(s) = \arg \max_{a \in \mathcal{A}} \mathcal{R}_s^a + P_{ss'}^a V(s')$$

- ✓ Greedy policy improvement over $Q(s, a)$ is model-free

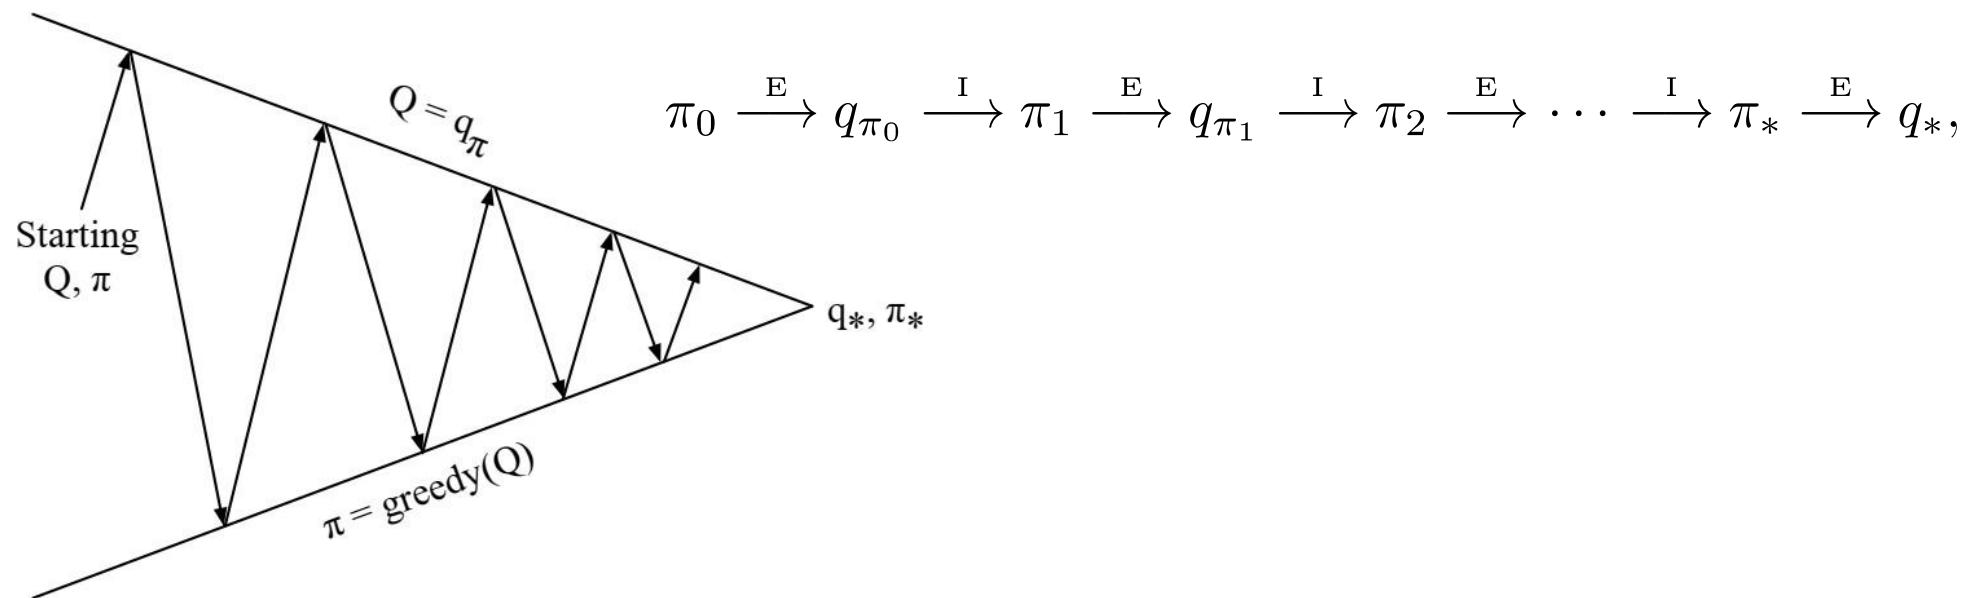
$$\pi'(s) = \arg \max_{a \in \mathcal{A}} Q(s, a)$$

Evaluating Action- Value Function with MC

$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a]$$

- ✓ REMINDER: First-visit and every-visit variants of MC
- ✓ Take averages for state-action pairs

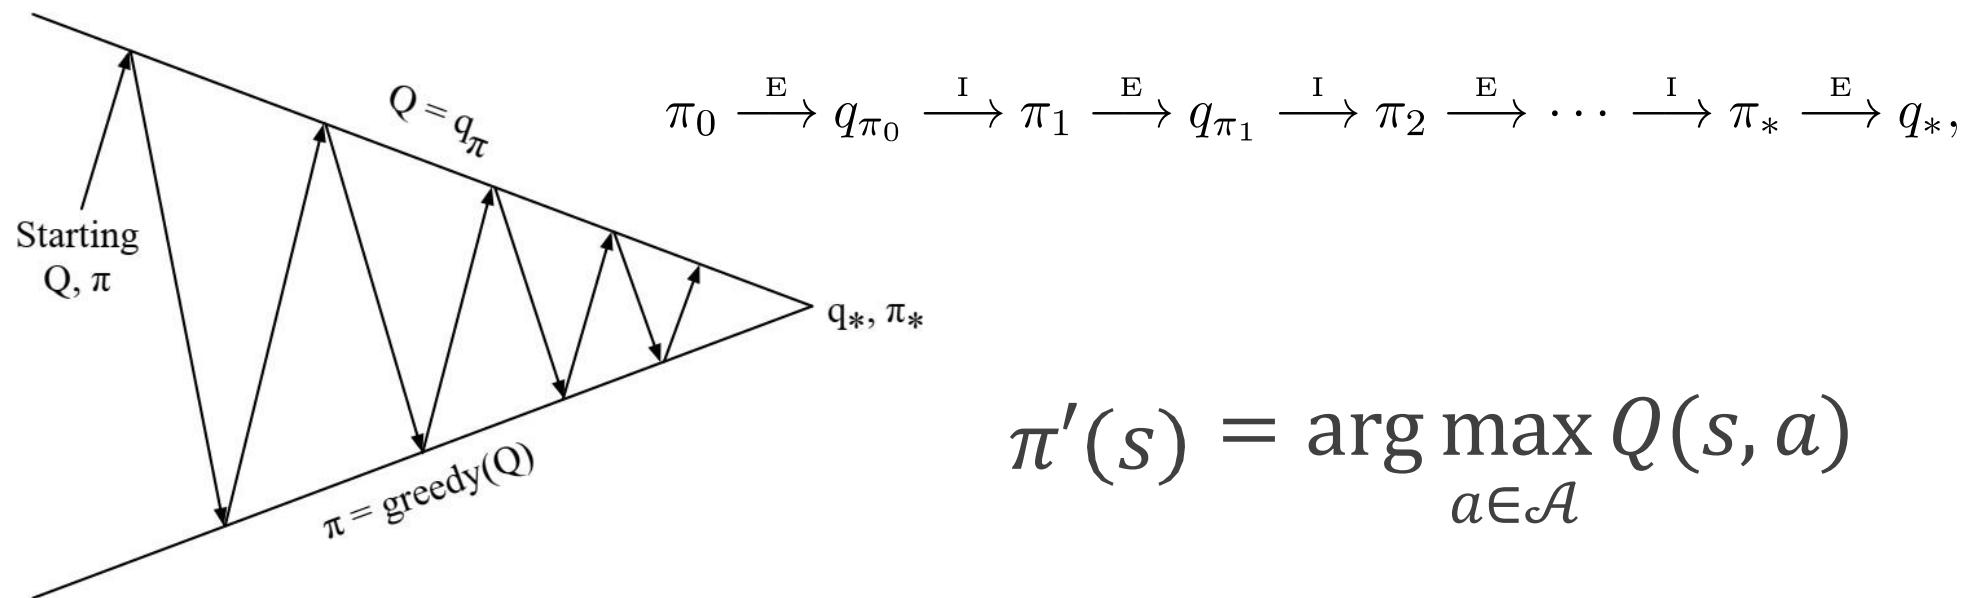
Generalized Policy Iteration with Action-Value Function



✓ Policy evaluation - Monte-Carlo policy evaluation, $Q = q_\pi$

✓ Policy improvement - Generate Greedy policy improvement?

Generalized Policy Iteration with Action-Value Function



- ✓ Policy evaluation - Monte-Carlo policy evaluation, $Q = q_\pi$
- ✓ Policy improvement - Generate Greedy policy improvement?

Evaluating Action- Value Function with MC

$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a]$$

- ✓ REMINDER: First-visit and every-visit variants of MC
- ✓ Take averages for state-action pairs
- ✓ **What happens if π is deterministic?**

Two approaches to ensure exploration in MC-control

- ✓ Exploring Starts (enforce pair state-action visiting).
- ✓ Soft policy

Exploring Starts MC-control

Monte Carlo ES (Exploring Starts), for estimating $\pi \approx \pi_*$

Initialize:

$\pi(s) \in \mathcal{A}(s)$ (arbitrarily), for all $s \in \mathcal{S}$

$Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Loop forever (for each episode):

Choose $S_0 \in \mathcal{S}, A_0 \in \mathcal{A}(S_0)$ randomly such that all pairs have probability > 0

Generate an episode from S_0, A_0 , following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:

Append G to $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow$ average($Returns(S_t, A_t)$)

$\pi(S_t) \leftarrow \operatorname{argmax}_a Q(S_t, a)$

Exploring Starts MC-control

Monte Carlo ES (Exploring Starts), for estimating $\pi \approx \pi_*$

Initialize:

$\pi(s) \in \mathcal{A}(s)$ (arbitrarily), for all $s \in \mathcal{S}$

$Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

What happens if we
don't know how to
explore the starts?

Loop forever (for each episode):

Choose $S_0 \in \mathcal{S}, A_0 \in \mathcal{A}(S_0)$ randomly such that all pairs have probability > 0

Generate an episode from S_0, A_0 , following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:

Append G to $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow$ average($Returns(S_t, A_t)$)

$\pi(S_t) \leftarrow \operatorname{argmax}_a Q(S_t, a)$

ϵ -greedy Exploration

- ✓ Simplest idea for ensuring continual exploration
- ✓ All m actions are tried with non-zero probability
 - ✓ With probability $1 - \epsilon$ choose the greedy action
 - ✓ With probability ϵ choose an action at random

$$\pi(a|s) = \begin{cases} \epsilon/m + (1 - \epsilon) & \text{if } a^* = \arg \max_{a \in \mathcal{A}} Q(s, a) \\ \epsilon/m & \text{otherwise} \end{cases}$$

ϵ -greedy Policy Improvement

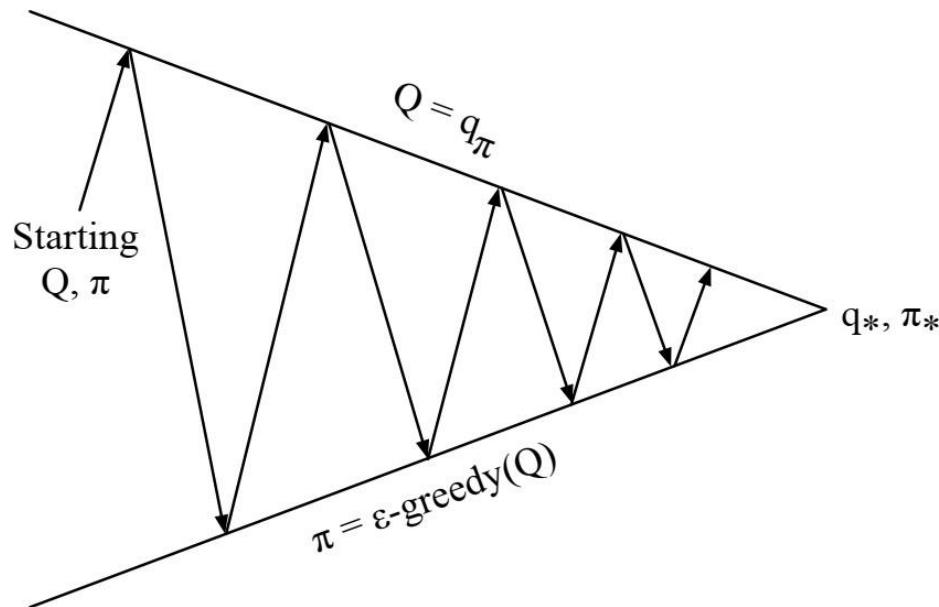
Theorem

For any ϵ -greedy policy π , the ϵ -greedy policy π' with respect to q_π is an improvement,
 $v_{\pi'}(s) \geq v_\pi(s)$

$$\begin{aligned} q_\pi(s, \pi'(s)) &= \sum_{a \in \mathcal{A}} \pi'(a|s) q_\pi(s, a) \\ &= \frac{\epsilon}{m} \sum_{a \in \mathcal{A}} q_\pi(s, a) + (1 - \epsilon) \max_{a \in \mathcal{A}} q_\pi(s, a) \\ &\geq \frac{\epsilon}{m} \sum_{a \in \mathcal{A}} q_\pi(s, a) + (1 - \epsilon) \sum_a \frac{\pi(a|s) - \epsilon/m}{1 - \epsilon} q_\pi(s, a) \\ &= \sum_{a \in \mathcal{A}} \pi(a|s) q_\pi(s, a) = v_\pi(s) \end{aligned}$$

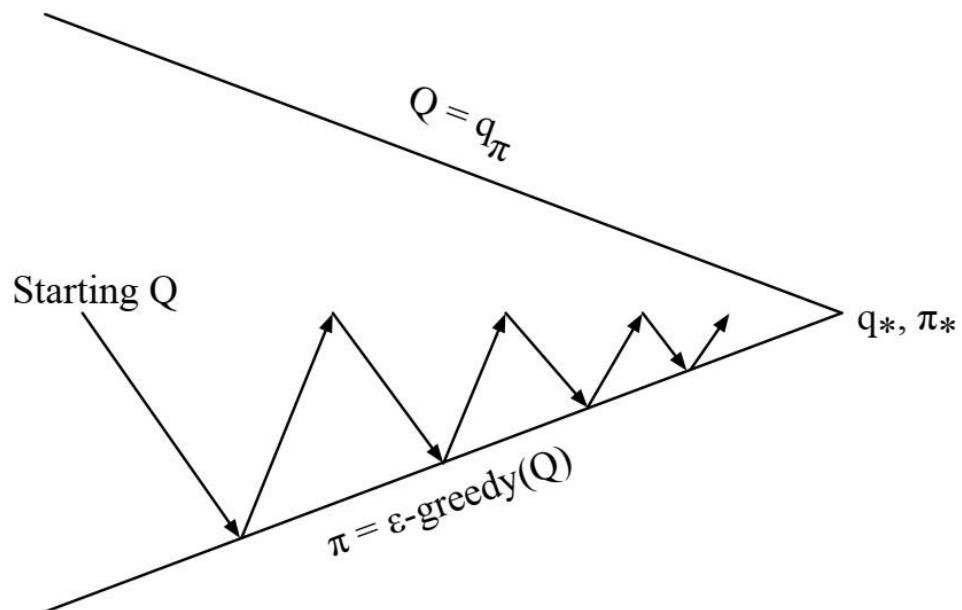
Therefore from **policy improvement theorem**
 $v_{\pi'}(s) \geq v_\pi(s)$

Monte-Carlo Policy Iteration



- ✓ Policy evaluation - Monte-Carlo policy evaluation, $Q = q_\pi$
- ✓ Policy improvement - ϵ -greedy policy improvement

Monte-Carlo Control



Every Episode

- ✓ Policy evaluation - Monte-Carlo policy evaluation, $Q \approx q_\pi$
- ✓ Policy improvement - ϵ -greedy policy improvement

Monte-Carlo Control

On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

Initialize:

$\pi \leftarrow$ an arbitrary ε -soft policy

$Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:

Append G to $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow$ average($Returns(S_t, A_t)$)

$A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)

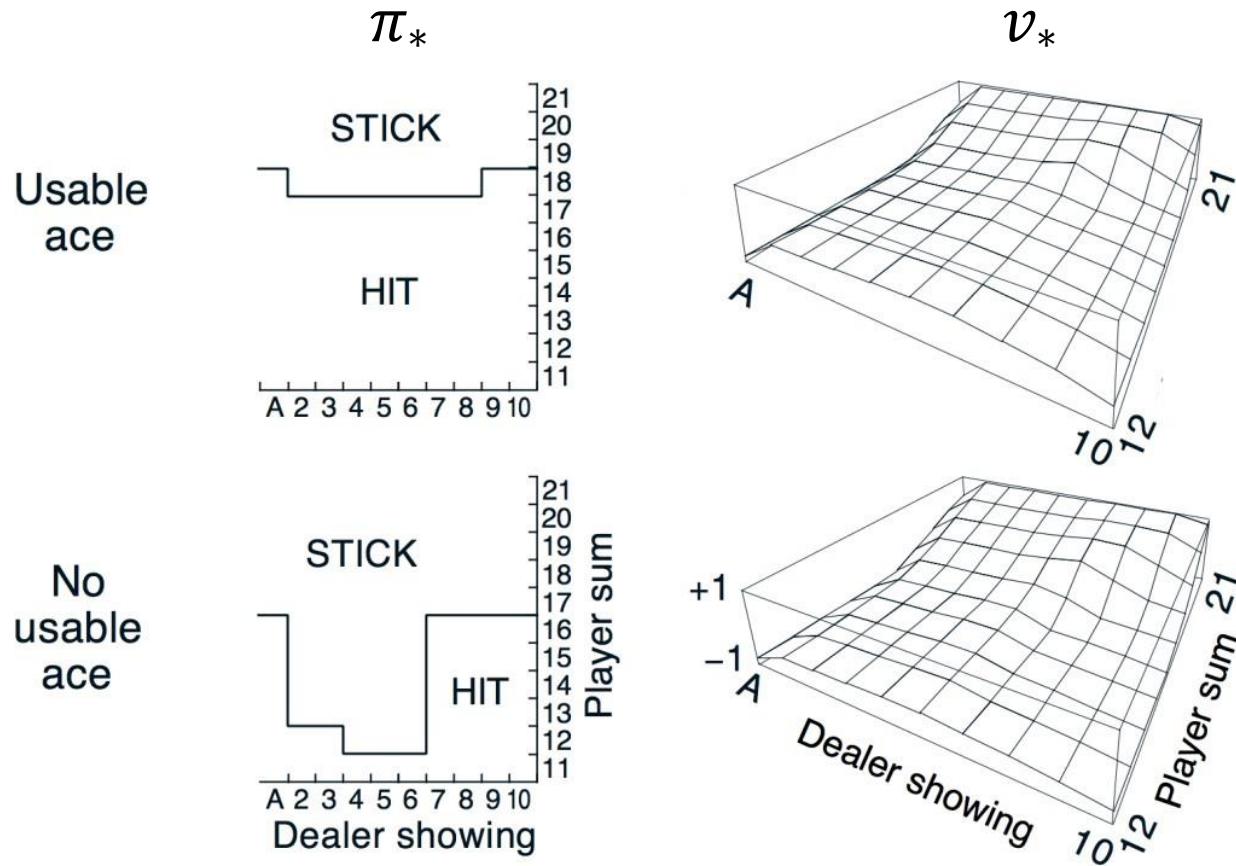
For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$



Blackjack Example

Blackjack – MC Control

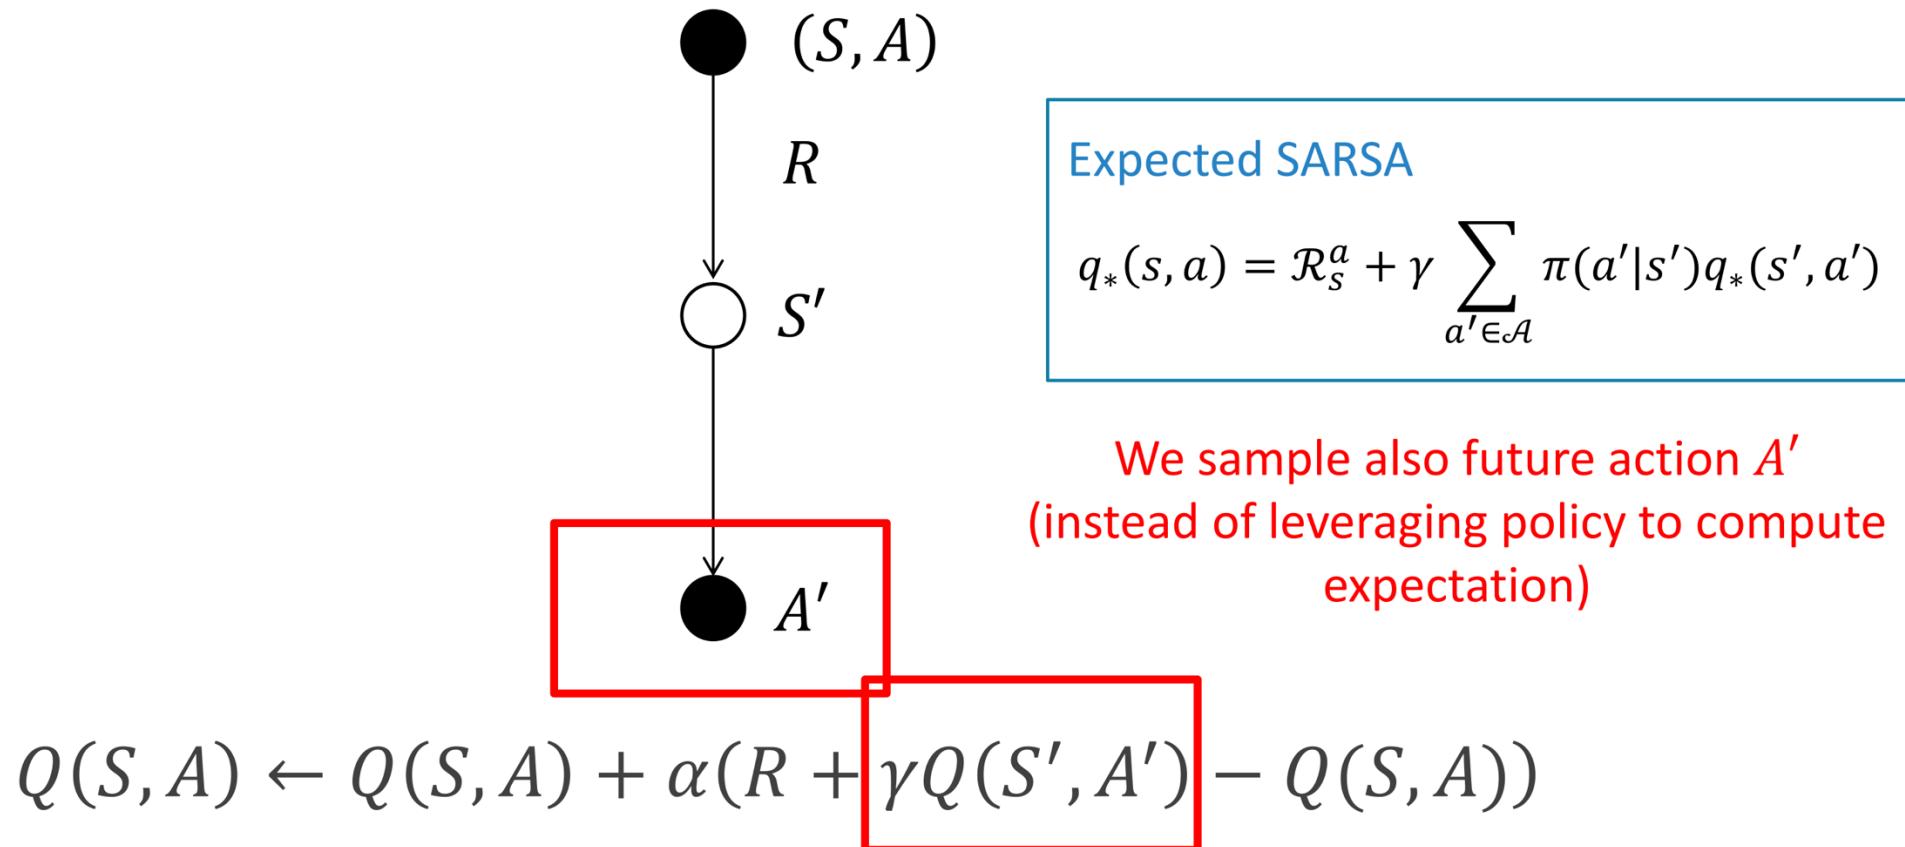


On-Policy TD Control

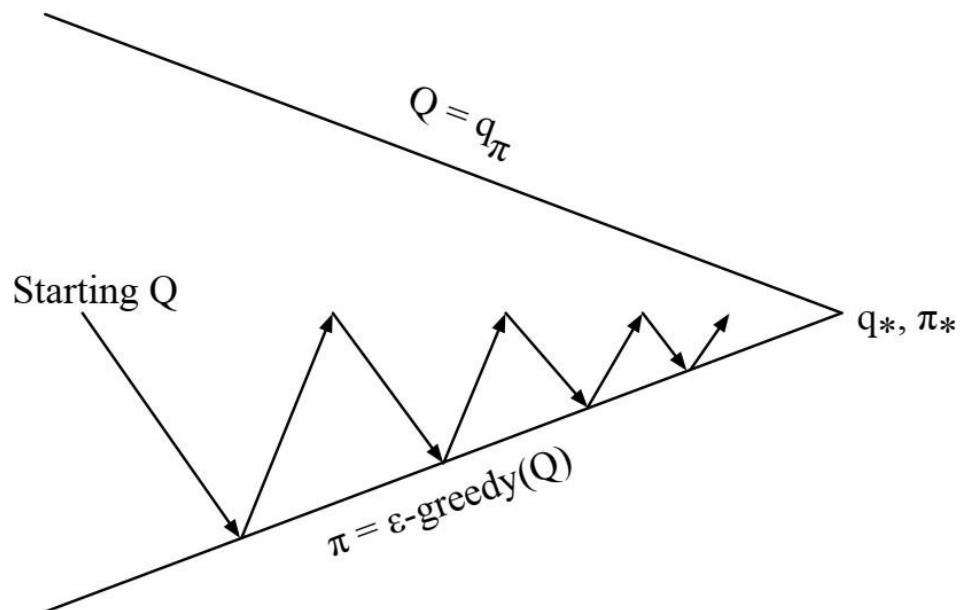
MC Vs TD Control

- ✓ TD learning has several advantages over MC
 - ✓ Lower variance
 - ✓ Online
 - ✓ Incomplete sequences
- ✓ Straightforward intuition - Use TD instead of MC in our control loop
 - ✓ Apply TD to $Q(s, a)$
 - ✓ Use ϵ -greedy policy improvement
 - ✓ Update every time-step

Updating Action-Value Functions with SARSA (State–action–reward–state–action)



On-Policy Control with SARSA



Every **time-step**

- ✓ Policy evaluation - **SARSA**, $Q \approx q_\pi$
- ✓ Policy improvement - ϵ -greedy policy improvement

SARSA Algorithm for On-Policy Control

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Loop for each step of episode:

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$$

$S \leftarrow S'; A \leftarrow A'$;

 until S is terminal

Time for TD Demo

https://cs.stanford.edu/people/karpathy/reinforcejs/gridworld_td.html

SARSA(λ)

n -step SARSA

- ✓ Consider the following n -step returns for $n = 1, 2, \dots, \infty$

$$n = 1 \quad (\text{SARSA}) \quad q_t^{(1)} = R_{t+1} + \gamma Q(S_{t+1}, A_{t+1})$$

$$n = 2 \quad q_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma Q(S_{t+2})$$

...

$$n = \infty \quad (\text{MC}) \quad q_t^{(\infty)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$$

- ✓ Define the n -step Q-return

$$q_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n Q(S_{t+n}, A_{t+n})$$

- ✓ n -step SARSA updates $Q(S, A)$ towards the n -step Q-return

$$Q(S, A) \leftarrow Q(S, A) + \alpha \left(q_t^{(n)} - Q(S, A) \right)$$

1-step Sarsa
aka Sarsa(0)



2-step Sarsa



3-step Sarsa



n-step Sarsa

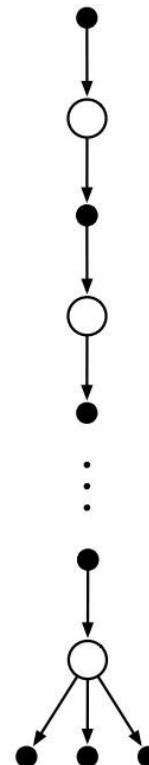
...



∞ -step Sarsa
aka Monte Carlo

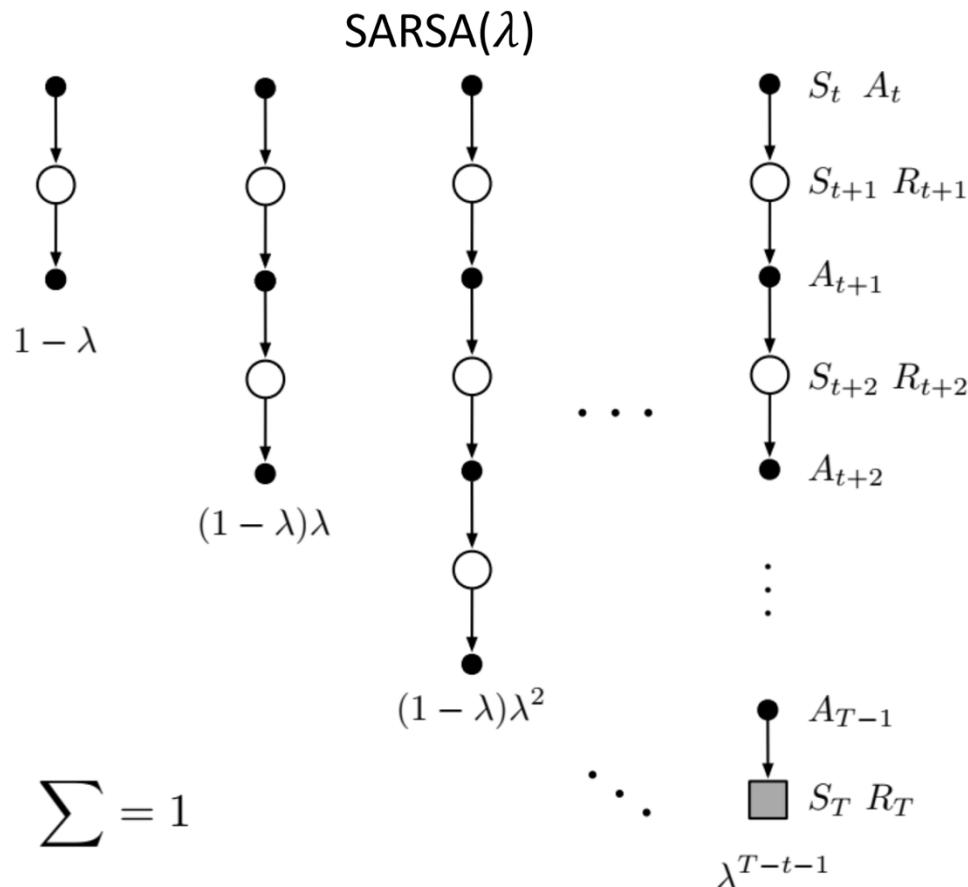


n-step
Expected Sarsa



SARSA
backups

SARSA(λ) - Forward View



✓ The q^λ return combines all n-step Q-returns $q_t^{(n)}$

✓ Using weight $(1 - \lambda)\lambda^{n-1}$

$$q_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} q_t^{(n)}$$

✓ Forward SARSA update

$$Q(S, A) \leftarrow Q(S, A) + \alpha(q_t^\lambda - Q(S, A))$$

SARSA(λ) - Backward View

- ✓ The return of eligibility traces
- ✓ SARSA(λ) needs one eligibility trace for each state-action pair

$$E_0(s, a) = 0$$
$$E_t(s, a) = \gamma \lambda E_{t-1}(s, a) + \mathbf{1}(S_t, A_t; s, a)$$

- ✓ $Q(s, a)$ is updated for every state s and action a in proportion to TD-error δ_t and eligibility trace $E_t(s, a)$

$$\delta_t = R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)$$
$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \delta_t E_t(s, a)$$

SARSA(λ) Algorithm

Initialize $Q(s, a)$ arbitrarily, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Repeat (for each episode):

$E(s, a) = 0$, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Initialize S, A

Repeat (for each step of episode):

Take action A , observe R, S'

Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$\delta \leftarrow R + \gamma Q(S', A') - Q(S, A)$

$E(S, A) \leftarrow E(S, A) + 1$

For all $s \in \mathcal{S}, a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow Q(s, a) + \alpha \delta E(s, a)$

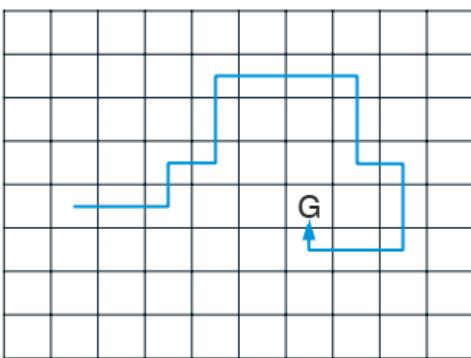
$E(s, a) \leftarrow \gamma \lambda E(s, a)$

$S \leftarrow S'; A \leftarrow A'$

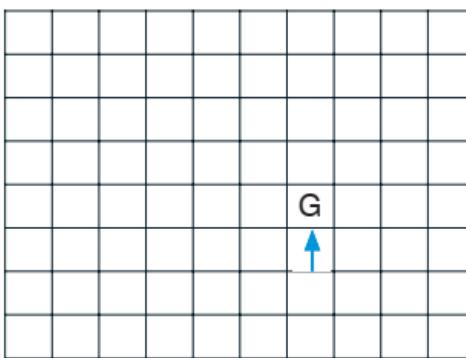
until S is terminal

SARSA(λ) on Gridworld

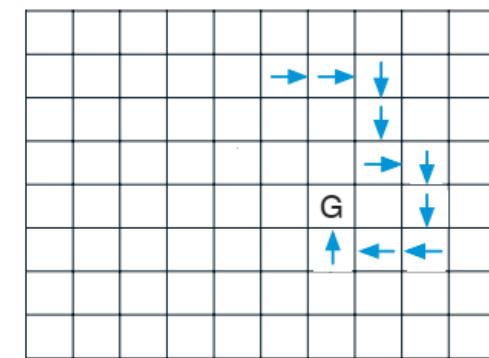
Path taken



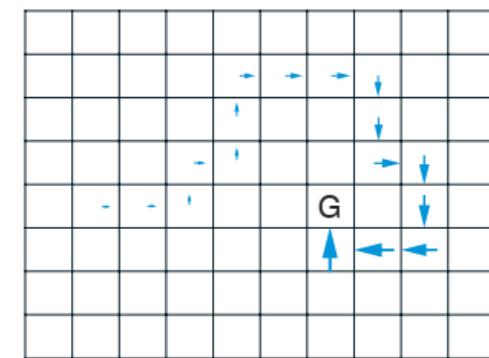
Action values increased
by one-step Sarsa



Action values increased by 10-step Sarsa



Action values increased by Sarsa(λ) with $\lambda=0.9$



Off-policy TD Learning

Off-Policy Learning

✓ Evaluate target policy $\pi(a|s)$ to compute $v_\pi(s)$ or $q_\pi(s, a)$

✓ While following behaviour policy $\mu(a|s)$

$$\{S_1, A_1, R_2, \dots, S_T\} \sim \mu$$

✓ Why is this important?

✓ Learn from imitation (humans, other agents,...)

✓ Re-use experience generated from old policies $\pi_1, \pi_2, \dots, \pi_{t-1}$

✓ Learn about optimal policy while following exploratory policy

✓ Learn about multiple policies while following one policy

Importance Sampling

Estimate the expectation leveraging an external (importance) distribution

Draw samples from importance distribution $Q(X)$ rather than from $P(X)$

$$\begin{aligned}\mathbb{E}_{X \sim P}[f(X)] &= \sum P(X)f(X) \\ &= \sum Q(X)\frac{P(X)}{Q(X)}f(X) \\ &= \mathbb{E}_{X \sim Q}\left[\frac{P(X)}{Q(X)}f(X)\right]\end{aligned}$$

Assign weights such that the empirical expectation (on $Q(X)$ samples) matches the expectation under $P(X)$

Importance Sampling for Off-Policy Monte Carlo

- ✓ Use returns generated from μ to evaluate π
- ✓ Weight return G_t according to similarity between policies
- ✓ Multiply importance sampling corrections along whole episode

$$G_t^{\pi/\mu} = \frac{\pi(A_t|S_t)}{\mu(A_t|S_t)} \frac{\pi(A_{t+1}|S_{t+1})}{\mu(A_{t+1}|S_{t+1})} \dots \frac{\pi(A_T|S_T)}{\mu(A_T|S_T)} G_t$$

- ✓ Update value towards corrected return

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t^{\pi/\mu} - V(S_t))$$

- ✓ Importance sampling can dramatically increase variance

Importance Sampling for Off-Policy TD

- ✓ Use TD targets generated from μ to evaluate π
- ✓ Weight TD targets $R + \gamma V(S')$ by importance sampling
- ✓ Only need a single importance sampling correction

$$V(S_t) \leftarrow V(S_t) + \alpha \left(\frac{\pi(A_t | S_t)}{\mu(A_t | S_t)} (R_{t+1} + \gamma V(S_{t+1}) - V(S_t)) \right)$$

- ✓ Much lower variance than MC
- ✓ Policies only need to be similar over a single step

Q-Learning

Off-policy learning of action-values $Q(s, a)$

- ✓ No importance sampling is required
 - ✓ Next action is chosen using behaviour policy $A_{t+1} \sim \mu(\cdot | S_t)$
 - ✓ But we consider alternative successor action $A' \sim \pi(\cdot | S_t)$
 - ✓ And update $Q(S_t, A_t)$ towards value of alternative action
- $$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A') - Q(S_t, A_t))$$

Off-policy Control by Q-Learning

✓ Allow both behaviour and target policies to improve

✓ The target policy π is greedy w.r.t. $Q(S_t, A_t)$

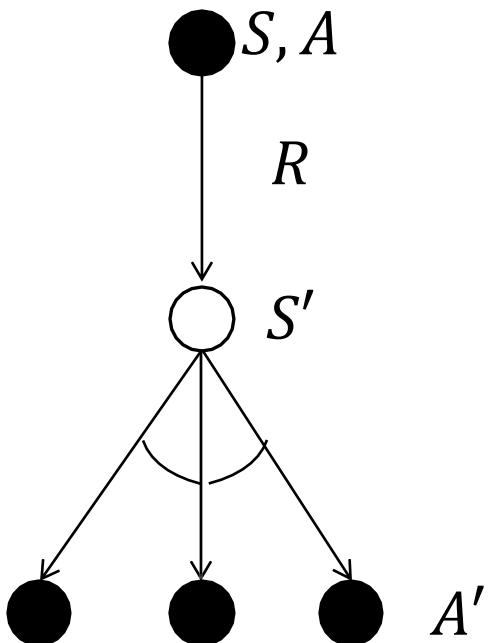
$$\pi(S_{t+1}) = \arg \max_a Q(S_{t+1}, a')$$

✓ The behaviour policy μ is ϵ -greedy w.r.t. $Q(s, a)$

✓ The Q-learning target then simplifies to

$$\begin{aligned} R_{t+1} + \gamma Q(S_{t+1}, A') &= R_{t+1} + \gamma Q\left(S_{t+1}, \arg \max_a Q(S_{t+1}, a')\right) \\ &= R_{t+1} + \max_a \gamma Q(S_{t+1}, a') \end{aligned}$$

Q-Learning Control Algorithm



Theorem

Q-learning control converges to the optimal action-value function, $Q(s, a) \rightarrow q_*(s, a)$

$$Q(S, A) \leftarrow Q(S, A) + \alpha \left(R + \max_a \gamma Q(S', a') - Q(S, A) \right)$$

Q-Learning Algorithm for Off-policy Control

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

 Initialize S

 Repeat (for each step of episode):

 Choose A from S using policy derived from Q (e.g., ε -greedy)

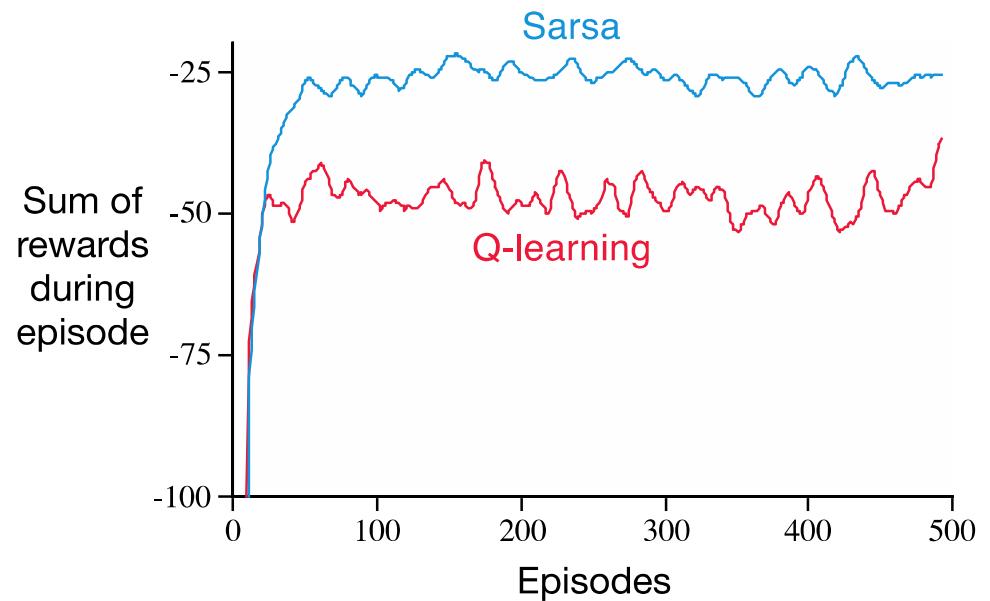
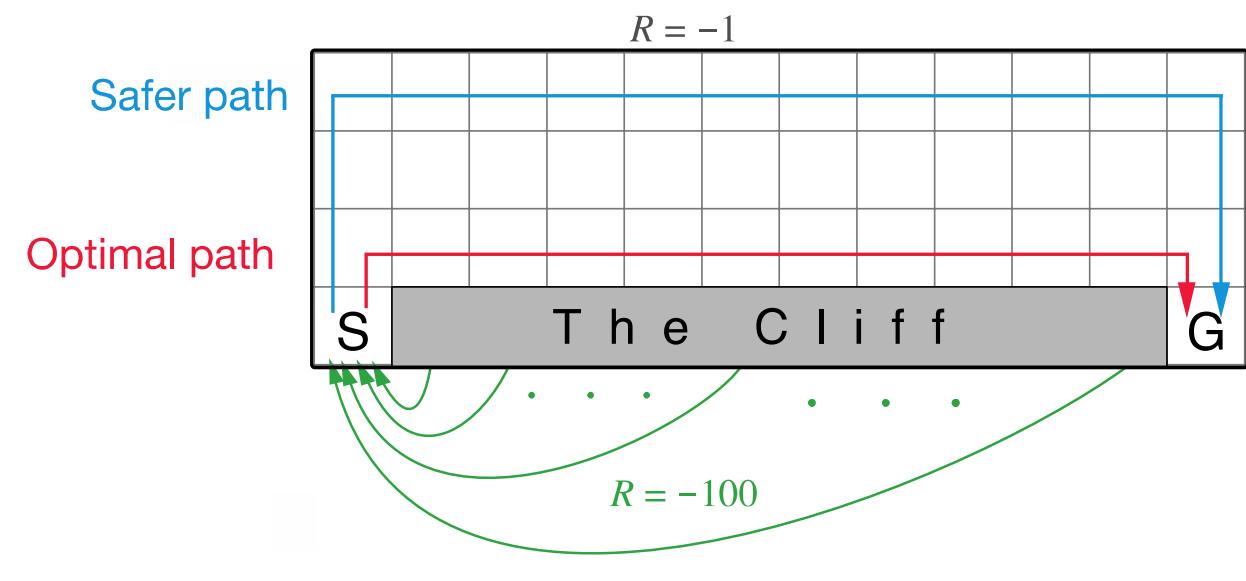
 Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$S \leftarrow S'$;

 until S is terminal

Q-Learning vs SARSA



OPTIMAL PATH vs ON-LINE PERFORMANCE

Q-learning & Exploration Demo

<https://www.aslanides.io/aixijs/demo.html>

Wrap-up

Dynamic Programming Vs Temporal Difference Learning

	<i>Full Backup (DP)</i>	<i>Sample Backup (TD)</i>
Bellman Expectation Equation for $v_\pi(s)$	<p>$v_\pi(s) \leftrightarrow s$</p> <p>Iterative Policy Evaluation</p>	<p>TD Learning</p>
Bellman Expectation Equation for $q_\pi(s, a)$	<p>$q_\pi(s, a) \leftrightarrow s, a$</p> <p>Q-Policy Iteration</p>	<p>Sarsa</p>
Bellman Optimality Equation for $q_*(s, a)$	<p>$q_*(s, a) \leftrightarrow s, a$</p> <p>Q-Value Iteration</p>	<p>Q-Learning</p>

Take (stay) home messages

- ✓ Model-Free control leverages action-value function
 - ✓ Greedy policy improvement does not need MDP
 - ✓ Generalized policy iteration
- ✓ Need to maintain sufficient exploration (ϵ -greedy)
- ✓ Off-policy control
 - ✓ Learning value function of a target policy from data generated by a different behaviour policy
 - ✓ Importance sampling to match the expectations of two policies
- ✓ TD control
 - ✓ On-policy: SARSA(λ)
 - ✓ Off-policy: Q-learning

Next Lecture

Value-function approximation

- ✓ Leave aside tabular environments
- ✓ Estimate value function with function approximation
- ✓ Linear models & neural networks
- ✓ MC & TD with Stochastic Gradient
- ✓ Experience replay buffers